

CONTENTS.

1	Introduction	1
2	The Transformer	3
2.1	Softmax	3
2.2	Embedding and Positional Encoding	4
2.3	The Attention Layer	4
2.4	Feed Forward Layer	6
2.5	Unembedding	7
2.6	Summary	8
2.7	Long-Context Transformers	8
3	The Need for Attention Scaling	8
3.1	Insufficiency of Standard Softmax	8
3.2	Scalable Softmax	10
4	Simplex-Like Geometry	12
4.1	Setup	12
4.2	Critical Scaling	13
4.3	Backpropagation	20
5	Attention at Initialization	26
5.1	Independent Entries	26
5.2	Correlated Entries	27
6	Deterministic Score Theory	34
6.1	Scaling Order	34
6.2	Rank-Collapse Boundary	37
6.3	Resolved Accumulation	39
	References	41

1 INTRODUCTION

At the most basic level, a natural language model is a function that takes as input a sequence of words and attempts to predict the word that follows. This prediction takes the form of a probability distribution over the set of all possible words, indicating the level of confidence the model has that any given word might follow. This model generates text on a given input by sampling from the output distribution. Once a word is generated, it is concatenated to the end of the original input, leading to a new input for the model to once again sample from. Repeating this process yields output text as desired.

More precisely, text is segmented into *tokens*. Although tokens are often words, they take on many other forms: punctuation, prefixes and postfixes, single characters. Ideally, these tokens would allow for a model that is not restricted to one language or type of text, and could make meaningful predications on any form of structured text with sufficient training data (other languages, code, etc.). We will assume that the set $\mathcal{X}$ of tokens for our model is chosen beforehand, but note that determining this set is an important part of engineering an efficient language model.

```

A genius is the one most like himself.
- Thelonious Monk

Ma parole, c'est fait! 50 buts!

三人行，必有我师。

\frac{\bar{X}_n - \mu}{\sigma / \sqrt{n}}
\rightarrow \mathcal{N}(0,1)

while True:
    print("This will never end!")

```

Figure 1: GPT-4 tokenizer on sample text.

Let n be the number of tokens in the sequence of text the model takes in as input, known as *context length*. The model can therefore be seen as a function $f : \mathcal{X}^n \rightarrow \mathcal{P}_{\mathcal{X}}$, where $\mathcal{P}_{\mathcal{X}}$ is the set of probability distributions over $\mathcal{X}$. To develop a model, we design the model's parametric form $f(\mathbf{X} \mid \theta)$ and use likelihood estimation on a large set of sample data to determine parameters θ for which model performs sufficiently well. These are not simple tasks: for any decent model, f is extremely complex and highly nonlinear, and the size of $\mathcal{X}$, n , and θ are several orders of magnitude beyond the regime on which classical methods in statistics are effective.

Model	Tokens $\mathcal{X}$	Context n	Parameters θ	Release
BERT (Large)	30000	512	340 Million	2018
GPT-2	50257	1024	1.5 Billion	2019
GPT-3	50257	2048	175 Billion	2020
Llama 3	128000	128000	405 Billion	2024

Table 1: Size of $\mathcal{X}$, n , and θ of some modern language models.

Assuming that there exists methods to optimize parameters θ with respect to model performance, this approach to developing a model is inherently powerful. Models can first be designed using intuition about how a machine may interpret and predict language under a mathematical framework. If the model performs sufficiently well after training, the goal is accomplished. From a purely predictive perspective, it is not necessary that the learned internal mechanisms resemble the original intuition motivating the architecture.

Earlier approaches to language modeling relied on architectures (RNN, CNN) which primarily capture information in a limited or compressed way as text is processed sequentially. These models have difficulty maintaining a truly global understanding of a longer passage, since dependencies from distant parts of the text are either truncated or gradually lost as information is passed through successive steps.

Introduced in 2017, the Transformer architecture [Vas+17] resolves this issue by introducing a completely different mechanism, one which allows every token to directly interact with every other token in the input. This shift fundamentally changed the field of natural

language processing, ultimately powering the large language models that now underpin modern AI systems. Section 2 gives an intuition-based overview of the Transformer architecture.

Among other ways to scale a model, of particular interest is the context length n . In theory, this would allow the model to use much longer-range context, enabling it to connect information across distant parts of a document and maintain a more coherent global understanding of the input. In real-world applications, this is especially useful for tasks like summarizing long reports, analyzing legal or financial documents, and answering questions that require information spread across many pages. It also improves performance in settings like code generation, where earlier definitions or functions must be remembered much later in the file.

However, scaling n has been shown to be a non-trivial task. Among other reasons, the attention mechanism core to the Transformer is susceptible to degenerate behaviour in the large- n limit, which leads to complete loss of expressive power of the model. Recent work has focused on improving the scaling behaviour of softmax attention by introducing an extra parameter, which can be found proposed in [Nak25]. Empirically, various modifications have been shown to mitigate the effects of this collapsing attention (REFERENCES), but mathematical justification has remained incomplete. The rest of this report summarizes relevant results in an attempt to unify existing empirical observations and theoretical perspectives on attention scaling.

2 THE TRANSFORMER

This section provides a precise description of the architecture of the transformer model as outlined in [Vas+17]. For each step, a brief heuristic justification of the design choice is given based on intuition about how a machine might use the architecture to model language effectively, however it is important to keep in mind that all that is being done is designing the parametric form $f_\theta(\mathbf{X})$, from which the model will optimize θ over a set of training data. Whether the model “thinks” this way in practice is hard to determine, and likely not the case.

Throughout, any *weight matrix* W or *bias vector* b is a learned parameter of the model, and any other value is either an input or an intermediary value. In a computational spirit, we will use the “ $\leftarrow$ ” symbol to indicate variable assignment. Vector arrows will be used in order to clarify when individual rows or columns of matrices are in question.

2.1 Softmax

Beyond language modelling, it is frequently needed that a model output a probability distribution over a finite set of candidate outcomes. This arises in tasks such as classification, retrieval, and decision-making, where the model must assign relative likelihoods to discrete alternatives rather than produce a single deterministic value. In these settings, the outputs are typically unconstrained real-valued scores that must be converted into a valid probability distribution.

Definition 2.1.1 (Softmax): Let $s \in \mathbb{R}^n$. Define the function $\text{softmax} : \mathbb{R}^n \rightarrow \mathcal{P}_n$ as

$$\text{softmax}(s)_j = \frac{e^{s_j}}{\sum_{k=1}^n e^{s_k}}.$$

For a matrix $S \in \mathbb{R}^{n \times m}$, $\text{softmax}(S)$ is the $\mathbb{R}^{n \times m}$ matrix where softmax has been applied either row-wise or column-wise, depending on context.

It is clear that softmax outputs a valid probability distribution, is differentiable, and preserves the ordering of scores. While many other functions could in principle be used, softmax is the most widely adopted choice.

2.2 Embedding and Positional Encoding

We assign each token in $\mathcal{X}$ is assigned a unique number from 1 to $|\mathcal{X}|$ and represent it as one-hot vector¹. Thus, our input $X = (\vec{X}_1, \dots, \vec{X}_n) \in \mathcal{X}^n$ is represented as a $|\mathcal{X}| \times n$ matrix, where each column is a one-hot encoding of a token.

Modeling and learning are much more difficult in high-dimensional discrete spaces, as the tools of calculus are not available. Hence, we choose an *embedding dimension* d which allows for a continuous representation of token “meaning” in $\mathbb{R}^d$. One interpretation motivating this embedding space is the idea that orthogonal directions could encode different concepts or ideas, and that linear combinations of vectors along these directions could represent the meaning of a given token.

A matrix $W_E \in \mathbb{R}^{d \times |\mathcal{X}|}$ contains all token embeddings, where column i is the embedding for token i . For our input X , the columns of the matrix multiplication $W_E X$ give the corresponding embeddings for each token. We also allow for the model to obtain information of the position of a token in the sequence. This is given by the matrix W_P , where column i contains the encoding for position i . This is added to the token embeddings to obtain

$$E \leftarrow W_E X + W_P,$$

which contains all token-level information that is independent of any interactions between tokens.

2.3 The Attention Layer

Given that words in any language rarely have a unique meaning independent of context, it is essential to allow tokens to “communicate” with each other. The attention layer modifies the positions of tokens in embedding space to better suit their meaning in context.

2.3.1 Single-Headed Attention

¹A one-hot encoding is a vector that has a 1 in the position corresponding to the number it represents, and 0 everywhere else.

Let $W_K, W_Q \in \mathbb{R}^{d_h \times d}$, where d_h is known as the *head dimension*. Given the previously obtained token embeddings E , we compute

$$Q \leftarrow W_Q E, \quad K \leftarrow W_K E,$$

called the *query matrix* and *key matrix*, respectively. Note that the columns of these matrices are transformations of embedding vectors: $\vec{Q}_i = W_Q \vec{E}_i$ and $\vec{K}_i = W_K \vec{E}_i$. The query $\vec{Q}_i$ represents what information the token is looking for from other tokens in order to refine its contextual meaning. The key $\vec{K}_i$ represents what kinds of information the token can provide to other tokens. Both these vectors are projections of their corresponding token embeddings into a shared lower-dimensional space $\mathbb{R}^{d_h}$, specifically used for relational matching.

To capture the extent to which the key $\vec{K}_j$ is relevant to query $\vec{Q}_i$, we compute the *attention score* $s_{ij} = \vec{Q}_i \cdot \vec{K}_j$, where larger values correspond to alignment in the query-key space. Thus, the matrix

$$S \leftarrow \frac{Q^T K}{\sqrt{d_h}}$$

has scaled attention score $s_{ij}/\sqrt{d_h}$ as entry (i, j) . This scaling factor keeps dot-product similarity “dimension-invariant”, meaning scores stay at a consistent scale regardless of head dimension.

Given a query, we would like these scores to correspond to weights indicating the relative importance of each key’s token for updating the embedding of the queried token. We therefore apply the softmax function column-wise to S giving

$$A \leftarrow \text{softmax}(S).$$

This completes the conversion of token embeddings into a normalized distribution over which tokens are most relevant to each other.

Returning to the original embeddings E , let $W_V \in \mathbb{R}^{d \times d}$. we compute the *value matrix*

$$V \leftarrow W_V E.$$

The columns $\vec{V}_i = W_V \vec{E}_i$ represent the information each token contributes when it is used by other tokens. Recalling that the columns $\vec{A}_i$ of the matrix A contain the attention weights assigned by token i , the update to a token’s embedding is given by the weighted sum of the value vectors according to weights $\vec{A}_i$. In matrix form, this can be written as

$$\Delta E \leftarrow V A.$$

This update is then applied to E to obtain the final output of the attention layer,

$$E_{\text{out}} \leftarrow E + \Delta E.$$

2.3.2 Multi-Headed Attention

We wish to use multiple single-headed attention layers in parallel to allow different heads to specialize in different types of relationships and patterns within the input sequence. Let h be the number of heads in this multi-headed attention layer. We choose the head dimension d_h such that $d = hd_h$.

As before, for each head, we have matrices $W_Q^{(i)}, W_K^{(i)} \in \mathbb{R}^{d_h \times d}$ and compute

$$A^{(i)} \leftarrow \text{softmax} \left(\frac{Q^T K}{\sqrt{d_h}} \right).$$

Allowing for each $W_V^{(i)}$ matrix to have size $d \times d$ would result in hd^2 parameters across all heads. In practice, we would like the parameter count to be comparable to W_Q and W_K , so we force W_V to have the low-rank structure

$$W_V^{(i)} = W_{V_l}^{(i)} W_{V_r}^{(i)},$$

where $W_{V_l}^{(i)} \in \mathbb{R}^{d \times d_h}$ and $W_{V_r}^{(i)} \in \mathbb{R}^{d_h \times d}$. Across all heads, this contributes $2dd_h$ parameters.

While this factored matrix could be applied per-head, each attention layer instead computes $W_{V_r}^{(i)} A$, which are then concatenated vertically and multiplied by the horizontal concatenation of $W_{V_l}^{(i)}$ matrices, giving²

$$\Delta E \leftarrow \underbrace{\begin{bmatrix} W_{V_l}^{(1)} & \dots & W_{V_l}^{(h)} \end{bmatrix}}_{W_O} \begin{bmatrix} W_{V_r}^{(1)} A \\ \vdots \\ W_{V_r}^{(h)} A \end{bmatrix},$$

which is equivalent to computing $\Delta E^{(i)} \leftarrow W_V^{(i)} A$ for every head and summing the results. Because of this, we will consider the horizontal concatenation W_O to be a separate matrix that is applied after concatenating the results from each individual head. For each head, we will consider the $W_{V_r}^{(i)}$ matrix to simply be the $W_V^{(i)}$ matrix applied in the single-headed case.

Adding this to the original embeddings gives the final output for the multi-headed attention,

$$E_{\text{out}} \leftarrow E + \Delta E.$$

2.3.3 Masking and LayerNorm

MASKING LAYERNORM

2.4 Feed Forward Layer

Once a token has been contextualized using attention, it carries information from across the sequence, but this information is still only combined linearly through a weighted sum. The

²This allows the transformation to be expressed as a composition of two learned linear maps, which is computationally more efficient and better structured for parameter optimization than a product of two matrices.

feed forward layer allows for each token to independently introduce nonlinear transformations in its embedding, converting this context into more expressive and higher-level representations. This is done using a *multi-layer perceptron*.

Let E be input embeddings that have been processed by an attention layer. The multi-layer perceptron aims to model how a biological brain might process information using neurons. We choose a feed-forward dimension d_f , which we will interpret to be the number of neurons in this layer.

Each neuron processes information using an affine transformation of the input data. This is then fed into a non-linear *activation function* σ which aims to model how a neuron might fire. We use $\sigma(x) = \max(0, x)$, known as ReLU, although there are many other options. The *activation* of neuron i for some embedding $\vec{E}_j$ is the scalar value given by

$$n_{ij} \leftarrow \sigma\left(\overline{W}_{\text{in}}^{(i)} \cdot \vec{E}_j + b_{\text{in}}^{(i)}\right).$$

Letting W_{in} to be the matrix with weights $\overline{W}_{\text{in}}^{(j)}$ as columns and b_{in} to be the column vector with components $b_{\text{in}}^{(i)}$, this can be expressed in matrix form as

$$N \leftarrow \sigma\left(W_{\text{in}}E + b_{\text{in}}\mathbf{1}_{d_f}^T\right),$$

where σ is applied component-wise, and $\mathbf{1}_{d_f} \in \mathbb{R}^{d_f}$ is the column vector of all ones.

Similar to attention, an update to each embedding is calculated using an affine transformation of the embedding's corresponding neuron activations. In parallel, this gives

$$\Delta E \leftarrow W_{\text{out}}N + b_{\text{out}},$$

which is added to the input to give

$$E_{\text{out}} = E + \Delta E.$$

2.5 Unembedding

In practice, embeddings are fed through many iterations of transformer blocks (a multi-headed attention layer followed by a feed forward layer), each with separate parameters. After this processing, we wish to convert each embedding $\vec{E}_i$ into a probability distribution over tokens $\mathcal{X}$, indicating the level of confidence that each token in the vocabulary could be the next token at position i .

This is done by applying the unembedding matrix $W_U \in \mathbb{R}^{|\mathcal{X}| \times d}$ to $\vec{E}_i$, followed by a softmax. In matrix form,

$$Y = \text{softmax}(W_U E + b_U \mathbf{1}^T),$$

where softmax is applied column-wise. Note that each column $\vec{Y}_i$ is the prediction for the next token at each position i in the input sequence, ignoring tokens that come after position i . This is useful for model training purposes, but for generation purposes we only use the distribution $\vec{Y}_n$ corresponding to the last input token.

2.6 Summary

2.7 Long-Context Transformers

While the term “long” continuously changes meaning as advances are made in machine learning, long-context transformers are transformers that are designed to support a context length n that is substantially longer than what is seen in common models. Many real-world tasks such as reasoning over long documents, codebases, or dialogues require access to information that may be introduced thousands or even millions of tokens earlier, where simply truncating the context leads to severely degraded performance. Increasing context length effectively augments the models “memory”, which allows for richer retrieval of relevant evidence and more coherent global reasoning.

It is important to distinguish between two types of long-context transformers. The first consists of models explicitly trained on long sequences, which attempt to internalize long-range dependencies during optimization. The second consists of models trained on shorter sequences but designed to generalize at inference time to longer sequences. There also exists hybrids of these two models, which might split training into 80% short sequences and 20% long sequences, for example.

From a scaling perspective, there are many issues related to naively increasing context length that apply to both types of long-context transformers. First and foremost, attention grows quadratically in context length, since a vanilla attention matrix A is of size $n \times n$. Second, attention suffers from large- n scaling issues inherent to the softmax function. The remainder of this report investigates the latter. In general, the results collected throughout are more so properties of the softmax function under specific inputs rather than anything regarding the underlying Transformer architecture or training, and thus we will not differentiate between the two types of long-context Transformers unless it is pertinent.

3 THE NEED FOR ATTENTION SCALING

3.1 Insufficiency of Standard Softmax

3.1.1 Single Head

Consider the standard softmax function applied to a sequence of score vectors $(s^{(n)})$ such that $s^{(n)} \in \mathbb{R}^n$ for all n with uniformly bounded components $m \leq s_j^{(n)} \leq M$ for $m, M \in \mathbb{R}$. For a fixed j ,

$$\text{softmax}(s^{(n)})_j = \frac{e^{s_j}}{\sum_{k=1}^n e^{s_k}} \leq \frac{e^M}{ne^m} = \frac{e^{M-m}}{n},$$

as well as

$$\text{softmax}(s^{(n)})_j \geq \frac{e^{m-M}}{n}.$$

Then, we have that $\text{softmax}(s^{(n)})_j = \Theta(\frac{1}{n})$ for all j and thus $\|\text{softmax}(s^{(n)})\|_\infty \rightarrow 0$ as $n \rightarrow \infty$. Heuristically, the entire softmax vector becomes diffuse and “true” attention, in the sense of assigning a non-vanishing proportion of mass to a distinguished set of coordinates, cannot emerge. Therefore, any nontrivial attention mechanism in the large- n limit must be driven by score fluctuations whose scale grows with n , allowing certain coordinates to overcome the $\Theta(n)$ growth of the softmax normalizing constant. Because attentions scores are usually designed to be probabilistically $O(1)$, this is naturally a problem when considering large- n inputs.

As an example, consider the sequence of vectors $s^{(n)} = (3, -2, \dots, -2)$ with first component 3 and all other components -2 , noting that these are clearly uniformly bounded. Figure 2 shows that not only does $a_{\max}$ decrease to 0 relatively quickly, but the entropy of the distribution quickly approaches the entropy of the uniform distribution. This is referred to as *rank collapse*, since it is the behaviour obtained when a rank one matrix is used in the attention calculation.




Figure 2: Rank collapse on scores with bounded components.

Even if scores are not uniformly bounded, the nature of the growth of the scores can lead to another failure mode known as *entropy collapse*. This occurs when the maximum component grows too quickly with respect to the rest of the scores, leading softmax to behave as a “hardmax” for large n , illustrated in Figure 3 for the sequence $t^{(n)} = \frac{1}{3} \log(n)s^{(n)}$.




Figure 3: Entropy collapse on scores with growing components.

3.1.2 Transformer Setting

Theorem 3.1.1 (Long-Context Attention Collapse):

Let $\mathcal{X} \subseteq \mathbb{R}^d$ be a finite set of token embeddings, and let $X^{(n)} \in \mathcal{X}^n$ be a matrix of n embeddings.

Define query and key matrices $Q^{(n)} = \varphi(X^{(n)}) \in \mathbb{R}^{n \times k}$, $K^{(n)} = \kappa(X^{(n)}) \in \mathbb{R}^{n \times k}$, where φ, κ are continuous functions obtained as compositions of a finite number L of layers of the following two types:

1. A feedforward layer $F_l(Z)_i = f_l(Z_i)$ that acts row-wise and is continuous,
2. Self-attention layer $G_l(Z)_i = \sum_{j=1}^n a_{ij} v_l(Z_j)$, where $a_{ij} \in [0, 1]$ are softmax-normalised attention coefficients and v is a continuous feedforward network.

Define the score matrix $S^{(n)} = Q^{(n)}(K^{(n)})^T \in \mathbb{R}^{n \times n}$ and row-wise softmax attention matrix $A^{(n)} = \text{softmax}(S^{(n)})$.

Then, $\|A^{(n)}\|_{\max} \rightarrow 0$ as $n \rightarrow \infty$, where $\|A^{(n)}\|_{\max} = \max_{ij} |a_{ij}|$.

[Vel+24]

Proof. Since $\mathcal{X} \subseteq \mathbb{R}^d$ is finite, both $\mathcal{X}$ and its convex hull $\text{conv}(\mathcal{X})$ are compact. Since all feedforward operators F_l and v_l act row-wise and are continuous, they preserve compactness row-wise. Each output row of a self-attention layer is a convex combination of vectors in a compact set, hence remains in a compact set. By induction over L , the finite number of layers, all embeddings lie in a compact set, independent of n . In particular, there exists a compact set $C \subseteq \mathbb{R}^k$ such that all rows $Q_i^{(n)}, K_i^{(n)} \in C$. Hence, there exists an $M > 0$ such that any score has the bound $|S_{ij}^{(n)}| \leq M$. By a similar derivation as before, we obtain

$$\frac{e^{-2M}}{n} \leq A_{ij}^{(n)} \leq \frac{e^{2M}}{n}.$$

Since M does not depend on n , $\|A^{(n)}\|_{\max} \rightarrow 0$ as $n \rightarrow \infty$. $\square$

While a simplified model, this shows a different fundamental limitation of Transformers. For a fixed trained model, there is no architectural mechanism that guarantees stable, non-degenerate attention behaviour as the sequence length n increases beyond the range seen during training. Although one might hope that the model could compensate by implicitly scaling scores with sequence length, the architecture does not provide a reliable way to represent or extrapolate n . Consequently, any implicit learned rescaling of W_K or W_Q matrices remains uniformly bounded, and cannot in general counteract the linear growth of the softmax normalization as $n \rightarrow \infty$.

3.2 Scalable Softmax

An initial instinct may be to introduce a fixed parameter $\beta > 0$ to softmax,

$$\text{softmax}(s)_j = \frac{e^{\beta s_j}}{\sum_{k=1}^n e^{\beta s_k}},$$

however this suffers the exact same problem as standard softmax when scores are bounded, since

$$\frac{e^{\beta(m-M)}}{n} \leq \text{softmax}(s)_j \leq \frac{e^{\beta(M-m)}}{n}.$$

Hence, any non-degenerate scaling β must depend on n , i.e. $\beta = \beta_n$. The *scalable softmax* function

$$\text{ssmax}(z)_j = \frac{e^{(sp_n+b)z_j}}{\sum_{k=1}^n e^{(sp_n+b)z_k}}$$

is proposed in [Nak25], where s and b are scalar learned parameters that are unique to each attention layer and head, and p_n is a learnable parameter shared across all layers and heads, depending only on the input vector size n . They train this scalable softmax on a standard model and find a relationship fitting $p_n \sim \log n$. For a given attention layer and head (i.e. fixed constants), this corresponds to $\beta_n \sim \log n$ as proposed in our original setup.

According to these results, consider the scaling $\beta_n = \gamma \log n$, for some $\gamma > 0$, and assume that s has a unique maximum $z_{\max} > z_{2\text{nd}}$. On one hand, we have

$$\|\text{softmax}(z)\|_{\infty} = \frac{n^{\gamma z_{\max}}}{\sum_{k=1}^n n^{\gamma z_k}} \leq \frac{n^{\gamma z_{\max}}}{(n-1)n^{\gamma z_{\min}} + n^{\gamma z_{\max}}} = \frac{1}{(n-1)n^{-\gamma(z_{\max}-z_{\min})} + 1}.$$

On the other hand,

$$\|\text{softmax}(z)\|_{\infty} \geq \frac{n^{\gamma z_{\max}}}{(n-1)n^{\gamma z_{2\text{nd}}} + n^{\gamma z_{\max}}} = \frac{1}{(n-1)n^{-\gamma(z_{\max}-z_{2\text{nd}})} + 1}.$$

If $\gamma > \frac{1}{z_{\max}-z_{\min}}$, we see that $\|\text{softmax}(z)\|_{\infty} \rightarrow 1$, whereas if $\gamma < \frac{1}{z_{\max}-z_{2\text{nd}}}$, then $\|\text{softmax}(z)\|_{\infty} \rightarrow 0$.

This is one example showing that even if β_n is of correct asymptotic order, constant multiples can still induce degenerate behaviour. Thus, β_n is quite fragile, which motivates the study of attention scaling beyond the “quick fix” of adding a learnable parameter to the model.

Returning to the example $s^{(n)} = (3, -2, -2, \dots, -2)$, we apply the scaling $\beta_n = 0.2 \log n$. Figure 4 shows that this scaling indeed appears to avoid the issues discussed above.




Figure 4: Non-degenerate behaviour of scaled softmax on $s^{(n)}$.

It is useful to rederive the critical scaling of this simple example as

$$\|\text{softmax}(z)\|_\infty = \frac{e^{3\beta}}{\underbrace{e^{3\beta}}_{\text{max}} + \underbrace{(n-1)e^{-2\beta}}_{\text{bulk}}} = (1 + e^{-5\beta + \log n} + o(e^{-5\beta + \log n}))^{-1},$$

which must have $\beta = 0.2 \log n$ in order to avoid convergence to either 0 or 1.

At this point, the question remains how to determine an effective scaling β_n in more general contexts, and whether any universality properties hold. Since transformer inputs lie between deterministic and fully random regimes, a completely general analysis seems intractable at the moment, and it is natural to first study simplified models that capture the essential structure of the attention mechanism before addressing more realistic settings.

4 SIMPLEX-LIKE GEOMETRY

The following is an overview of results and proofs found in [Che+25] with intermediate steps filled in.

4.1 Setup

Consider a simplified version of unmasked single-headed attention in which the head dimension equals the embedding dimension and $W_K = W_Q = W_V = I_d$. For a collection of token embeddings $x_1, \dots, x_n \in \mathbb{R}^d$, define the normalized vectors $u_i = x_i / \|x_i\|$. The scaled attention scores for the normalized vectors u_i are thus given by

$$s_{ij} = \beta \langle u_i, u_j \rangle,$$

where $\langle \cdot, \cdot \rangle$ denotes the standard inner product on $\mathbb{R}^d$, and β is a scaling factor to be chosen. The corresponding attention matrix A has entries

$$a_{ij} = \frac{e^{s_{ij}}}{Z_i}, \quad Z_i = \sum_{k=1}^n e^{s_{ik}}.$$

The main object at study is the nonlinear operator $\text{ATT} : \mathbb{R}^d \rightarrow \mathbb{R}^d$, defined as

$$\text{ATT}(u_i) = \sum_{j=1}^n a_{ij} u_j.$$

Finally, recalling that in transformer architectures, the attention layer computes an update to the original embeddings through a residual connection, we define the updated embeddings and their normalized versions

$$x'_i = \text{ATT}(u_i) + \alpha x_i, \quad u'_i = \frac{x'_i}{\|x'_i\|},$$

where $\alpha \geq 0$ is the strength of the residual connection. Geometrically, α interpolates between purely attention-driven dynamics and identity-preserving behaviour.

4.2 Critical Scaling

We are interested in how the attention scaling β affects the geometry of the token embeddings, namely how pairwise angles evolve under the attention update. Equivalently, this amounts to studying the relationship between $\langle u'_i, u'_j \rangle$ and $\langle u_i, u_j \rangle$.

Consider a further simplified setup in which $\langle u_i, u_j \rangle = p \in (0, 1)$ if $i \neq j$ and $\|u_i\|^2 = q$. The key insight is that the behaviour of attention is governed by the normalization factor Z_i , which simplifies to

$$Z_i = e^\beta + (n-1)e^{p\beta}$$

under these assumptions. Thus, the asymptotic behaviour of Z_i depends on the competition between the self-attention term and the aggregate contribution of attention from all other tokens. Under the scaling $\beta = \gamma \log n$,

$$e^\beta = n^\gamma, \quad (n-1)e^{p\beta} \asymp n^{1+p\gamma}.$$

Hence, the dominant contribution to Z_i depends on the comparison between the exponents γ and $1+p\gamma$. The critical threshold occurs when $\gamma = \frac{1}{1-p}$. This reveals that there are three regimes under which attention operates. In the subcritical regime $\gamma < \frac{1}{1-p}$, attention weights become asymptotically uniform, so each token moves toward the global average. In the supercritical regime $\gamma > \frac{1}{1-p}$, the self-attention term dominates, and attention operator converges to the identity map. In the critical regime $\gamma = \frac{1}{1-p}$, both terms are balanced, leading to nontrivial dynamics as desired for an effective attention layer.

In Theorem 4.2.1 we generalize this setup beyond the symmetric simplex setting by allowing pairwise inner products to vary within fixed bounds. We show that the attention operator acts asymptotically as a contraction in the subcritical regime and as an identity in the supercritical regime. Corollary 4.2.1 provides exact results for the limiting inner product for the simplex case. Although [Che+25] claims that these results demonstrate that the scaling $\beta = \Theta(\log n)$ is intrinsic to the behaviour of attention mechanisms rather than an artifact of the simplex geometry, we will later see that this in fact not the case.

Theorem 4.2.1 (Phase Transition in Attention Geometry):

Assume there exist constants $q_1, q_2 > 0$ and $p_1, p_2 \in (0, 1)$ such that $q_1 \leq \|x_i\|^2 \leq q_2$ and $p_1 \leq \langle u_i, u_j \rangle \leq p_2$ for any $i \neq j$, with $p_1 = \langle u_i, u_j \rangle$ for some i, j . Let $\beta = \gamma \log n$ with $\gamma > 0$.

1. If $\gamma < \frac{1}{1-p_1}$ then there is a constant $\varepsilon > 0$ depending on α, p_2, q_1, q_2 such that

$$\liminf_{n \rightarrow \infty} \min_{i \neq j} \langle u'_i, u'_j \rangle \geq p_1 + \varepsilon.$$

2. If $\gamma > \frac{1}{1-p_2}$, then for any i ,

$$\lim_{n \rightarrow \infty} \langle u'_i, u'_j \rangle = \langle u_i, u_j \rangle.$$

[Che+25]

In the subcritical regime, the value of ε could in theory be very small, and thus the theorem should be interpreted as a qualitative phase transition rather than a quantitative improvement guarantee in this regime. In the supercritical regime, we obtain a much stronger result. Although the randomly sampled vectors found in Figure 5 do not necessarily satisfy the assumptions of Theorem 4.2.1, we still observe a clear separation between the subcritical (left) and supercritical (right) regimes, consistent with the qualitative behaviour described by the theorem.




Figure 5: Behaviour of geometric attention on random vectors.

We will first prove a series of lemmas before proving Theorem 4.2.1.

Lemma 4.2.2: For any $i \in \{1, \dots, n\}$, we have

$$Z_i = \begin{cases} (1 + o_n(1)) \cdot \left(\sum_{k \neq i} e^{s_{ik}} \right) & \gamma < \frac{1}{1-p_1}, \\ (1 + o_n(1)) \cdot e^\beta & \gamma > \frac{1}{1-p_2}. \end{cases}$$

Proof. Recall $s_{ik} = \beta \langle u_i, u_k \rangle$ and $\beta = \gamma \log n$, giving

$$Z_i = \sum_{k=1}^n e^{s_{ik}} = n^\gamma + \sum_{k \neq i} n^{\gamma \langle u_i, u_j \rangle}$$

If $\gamma < \frac{1}{1-p_1}$, then $\gamma - \gamma p_1 - 1 < 0$. Using $\langle u_i, u_j \rangle \geq p_1$,

$$\frac{n^\gamma}{\sum_{k \neq i} n^{\gamma \langle u_i, u_j \rangle}} \leq \frac{n^\gamma}{(n-1)n^{\gamma p_1}} = n^{\gamma - \gamma p_1 - 1} \cdot \frac{1}{1 - 1/n}$$

which goes to 0 as $n \rightarrow \infty$. Hence,

$$Z_i = \left(1 + \frac{n^\gamma}{\sum_{k \neq i} n^{\gamma \langle u_i, u_j \rangle}} \right) \left(\sum_{k \neq i} n^{\gamma \langle u_i, u_j \rangle} \right) = (1 + o_n(1)) \left(\sum_{k \neq i} e^{s_{ik}} \right).$$

Similarly, if $\gamma > \frac{1}{1-p_2}$, then $\gamma p_2 + 1 - \gamma < 0$. Using $\langle u_i, u_j \rangle \leq p_2$,

$$\frac{\sum_{k \neq i} n^{\gamma \langle u_i, u_j \rangle}}{n^\gamma} \leq \frac{(n-1)n^{\gamma p_2}}{n^\gamma} = n^{\gamma p_2 + 1 - \gamma} (1 - 1/n)$$

which goes to 0 as $n \rightarrow \infty$. Hence,

$$Z_i = \left(1 + \frac{\sum_{k \neq i} n^{\gamma \langle u_i, u_j \rangle}}{n^\gamma} \right) (n^\gamma) = (1 + o_n(1)) \cdot e^\beta.$$

□

Lemma 4.2.3: If $\gamma > \frac{1}{1-p_2}$, then for any $i \in \{1, \dots, n\}$,

$$\text{ATT}(u_i) = u_i + o_n(1),$$

where $o_n(1)$ does not depend on i .

Proof. By the previous lemma,

$$\text{ATT}(u_i) = Z_i^{-1} \left(e^\beta u_i + \sum_{j \neq i} e^{s_{ij}} u_j \right) = (1 + o_n(1)) \left(u_i + e^{-\beta} \sum_{j \neq i} e^{s_{ij}} u_j \right),$$

using the fact that $(1 + o_n(1))^{-1} = 1 + o_n(1)$. Since $\|u_j\| = 1$ for all j ,

$$\left\| e^{-\beta} \sum_{j \neq i} e^{s_{ij}} u_j \right\| \leq e^{-\beta} \sum_{j \neq i} e^{s_{ij}} \leq n^{-\gamma} \cdot n^{\gamma p_2} (n-1)$$

which goes to 0 as $n \rightarrow \infty$, independent of i . □

Lemma 4.2.4: For any $i \in \{1, \dots, n\}$,

1. If $\gamma < \frac{1}{1-p_1}$,

$$\|x'_i\|^2 \leq \alpha^2 \|x_i\|^2 + 2\alpha \|x_i\| p_2 + p_2 + o_n(1).$$

2. If $\gamma > \frac{1}{1-p_2}$,

$$\|x'_i\|^2 = (\alpha \|x_i\| + 1)^2 + o_n(1).$$

In both cases, $o_n(1)$ does not depend on i .

Proof. Using the definition of the update rule, expanding the inner product gives

$$\|x'_i\|^2 = \alpha^2 \|x_i\|^2 + 2\alpha \|x_i\| \langle u_i, \text{ATT}(u_i) \rangle + \|\text{ATT}(u_i)\|^2.$$

For the case of $\gamma < \frac{1}{1-p_1}$, Lemma 3.1 gave $Z_i = (1 + o_n(1)) \left(\sum_{j \neq i} e^{s_{ij}} \right)$. Hence,

$$\begin{aligned} \langle u_i, \text{ATT}(u_i) \rangle &= \langle u_i, \frac{1}{Z_i} \sum_{j=1}^n e^{s_{ij}} u_j \rangle \\ &= \frac{1}{Z_i} \sum_{j=1}^n e^{s_{ij}} \langle u_i, u_j \rangle \\ &\leq \frac{1}{Z_i} \left(e^\beta + p_2 \sum_{j \neq i} e^{s_{ij}} \right) \\ &= p_2 + o_n(1), \end{aligned}$$

since $e^\beta / \sum_{j \neq i} e^{s_{ij}} = o_n(1)$, as shown previously. Similarly,

$$\begin{aligned} \|\text{ATT}(u_i)\|^2 &= \left\langle \frac{1}{Z_i} \sum_{k=1}^n e^{s_{ik}} u_k, \frac{1}{Z_i} \sum_{l=1}^n e^{s_{il}} u_l \right\rangle \\ &= \frac{1}{Z_i^2} \sum_{k=1}^n \sum_{l=1}^n e^{s_{ik} + s_{il}} \langle u_k, u_l \rangle \\ &\leq \frac{1}{Z_i^2} \left(e^{2\beta} + 2p_2 e^\beta \sum_{j \neq i} e^{s_{ij}} + p_2 \sum_{k \neq i} \sum_{l \neq i} e^{s_{ik} + s_{il}} \right) \\ &= p_2 + o_n(1). \end{aligned}$$

For the case of $\gamma > \frac{1}{1-p_2}$, applying Lemma 4.2.3 directly gives

$$\begin{aligned}
\|x'_i\|^2 &= \alpha^2 \|x_i\|^2 + 2\alpha \|x_i\| \langle u_i, u_i + o_n(1) \rangle + \langle u_i + o_n(1), u_i + o_n(1) \rangle \\
&= \alpha^2 \|x_i\|^2 + 2\alpha \|x_i\| + 1 + o_n(1) \\
&= (\alpha \|x_i\| + 1)^2 + o_n(1).
\end{aligned}$$

□

Lemma 4.2.5: For $i, j \in \{1, \dots, n\}$ with $i \neq j$,

1. If $\gamma < \frac{1}{1-p_1}$,

$$\langle x'_i, x'_j \rangle \geq p_1(\alpha \|x_i\| + 1)(\alpha \|x_j\| + 1) + o_n(1).$$

2. If $\gamma > \frac{1}{1-p_2}$,

$$\langle x'_i, x'_j \rangle = (\alpha \|x_i\| + 1)(\alpha \|x_j\| + 1) \langle u_i, u_j \rangle + o_n(1).$$

In both cases, $o_n(1)$ does not depend on i .

Proof. Expand the inner product to give

$$\begin{aligned}
\langle x'_i, x'_j \rangle &= \alpha^2 \|x_i\| \|x_j\| \langle u_i, u_j \rangle + \alpha \|x_i\| \langle u_i, \text{ATT}(u_j) \rangle \\
&\quad + \alpha \|x_j\| \langle u_j, \text{ATT}(u_i) \rangle + \langle \text{ATT}(u_i), \text{ATT}(u_j) \rangle.
\end{aligned}$$

If $\gamma < \frac{1}{1-p_1}$, by following a similar procedure as in Lemma 4.2.4 we can instead bound the terms in question from below, i.e.

$$\begin{aligned}
\langle u_i, \text{ATT}(u_j) \rangle &= \langle u_i, \frac{1}{Z_i} \sum_{k=1}^n e^{s_{jk}} u_k \rangle \\
&\geq \frac{1}{Z_i} \left(p_1 e^\beta + p_1 \sum_{j \neq i} e^{s_{ij}} \right) \\
&= p_1 + o_n(1),
\end{aligned}$$

and

$$\begin{aligned}
\langle \text{ATT}(u_i), \text{ATT}(u_j) \rangle &= \left\langle \frac{1}{Z_i} \sum_{k=1}^n e^{s_{ik}} u_k, \frac{1}{Z_j} \sum_{l=1}^n e^{s_{jl}} u_l \right\rangle \\
&\geq \frac{1}{Z_i Z_j} \left(p_1 e^{2\beta} + p_1 e^\beta \sum_{k \neq i} e^{s_{ik}} + p_1 e^\beta \sum_{l \neq j} e^{s_{jl}} + p_1 \sum_{k \neq i} \sum_{l \neq j} e^{s_{ik} + s_{jl}} \right) \\
&= p_1 + o_n(1).
\end{aligned}$$

Hence,

$$\begin{aligned}
\langle x'_i, x'_j \rangle &\geq \alpha^2 \|x_i\| \|x_j\| p_1 + \alpha (\|x_i\| + \|x_j\|) p_1 + p_1 + o_n(1) \\
&= p_1 (\alpha \|x_i\| + 1) (\alpha \|x_j\| + 1) + o_n(1).
\end{aligned}$$

If $\gamma > \frac{1}{1-p_2}$, Lemma 4.2.3 can be applied to obtain

$$\begin{aligned}\langle x'_i, x'_j \rangle &= \alpha^2 \|x_i\| \|x_j\| \langle u_i, u_j \rangle + \alpha \|x_i\| \langle u_i, u_j + o_n(1) \rangle \\ &\quad + \alpha \|x_j\| \langle u_j, u_i + o_n(1) \rangle + \langle u_i + o_n(1), u_j + o_n(1) \rangle \\ &= (\alpha \|x_i\| + 1)(\alpha \|x_j\| + 1) \langle u_i, u_j \rangle + o_n(1).\end{aligned}$$

□

We now prove Theorem 4.2.1.

Proof of Theorem 4.2.1. First, consider the case where $\gamma < \frac{1}{1-p_1}$. We can further simplify the result in Lemma 4.2.4 by completing the square, giving

$$\begin{aligned}\|x'_i\|^2 &\leq \alpha^2 \|x_i\|^2 + 2\alpha \|x_i\| p_2 + p_2 + o_n(1) \\ &= \alpha^2 \|x_i\|^2 + 2\alpha \|x_i\| + 1 - 2\alpha \|x_i\| - 1 + 2\alpha \|x_i\| p_2 + p_2 + o_n(1) \\ &= (\alpha \|x_i\| + 1)^2 - (1 - p_2)(2\alpha \|x_i\| + 1) + o_n(1) \\ &\leq (\alpha \|x_i\| + 1)^2 - (1 - p_2)(2\alpha \sqrt{q_1} + 1) + o_n(1),\end{aligned}$$

where the last inequality follows from the fact that $\|x_i\| \geq \sqrt{q_1}$. Thus, there is some constant $\delta > 0$ depending on α, p_2, q_1, q_2 such that

$$\frac{1}{\|x_i\|} \geq \frac{1 + \delta}{\alpha \|x_i\| + 1} + o_n(1).$$

An equivalent formulation of the result of Lemma 4.2.5 is the equation

$$\langle u'_i, u'_j \rangle \geq \frac{(\alpha \|x_i\| + 1)(\alpha \|x_j\| + 1)}{\|x_i\| \|x_j\|} p_1 + o_n(1),$$

and hence

$$\langle u'_i, u'_j \rangle \geq p_1(1 + \delta)^2 + o_n(1) = p_1 + \varepsilon + o_n(1),$$

where $\varepsilon = p_1(2\delta + \delta^2)$. Since this was for any $i \neq j$, it follows that

$$\liminf_{n \rightarrow \infty} \min_{i \neq j} \langle u'_i, u'_j \rangle \geq p_1 + \varepsilon.$$

For the case where $\gamma > \frac{1}{1-p_2}$, we formulate the result of Lemma 4.2.5 as

$$\langle u'_i, u'_j \rangle = \frac{(\alpha \|x_i\| + 1)(\alpha \|x_j\| + 1)}{\|x_i\| \|x_j\|} \langle u_i, u_j \rangle + o_n(1)$$

and substitute the result of Lemma 4.2.4 to immediately obtain

$$\langle u'_i, u'_j \rangle = \langle u_i, u_j \rangle + o_n(1).$$

Hence,

$$\lim_{n \rightarrow \infty} \langle u'_i, u'_j \rangle = \langle u_i, u_j \rangle.$$

□

Corollary 4.2.1 (Phase Transition, Simplex Case): Assume there exist constants $q > 0$, $p \in (0, 1)$ such that $\|x_i\|^2 = q$ and $\langle u_i, u_j \rangle = p$ for any $i \neq j$. Then,

$$\lim_{n \rightarrow \infty} \langle u'_i, u'_j \rangle = \begin{cases} \frac{p(\alpha\sqrt{q}+1)^2}{\alpha^2q+2\alpha\sqrt{q}p+p} & \gamma < \frac{1}{1-p}, \\ \frac{p(\alpha\sqrt{q}+1)^2}{\alpha^2q+\alpha\sqrt{q}(1+p)+\frac{1+3p}{4}} & \gamma = \frac{1}{1-p}, \\ p & \gamma > \frac{1}{1-p}. \end{cases}$$

Proof. This corresponds to the special case of Theorem 4.2.1 where $p_1 = p_2$.

If $\gamma < \frac{1}{1-p}$, the steps of the proof of Lemma 4.2.4 follow with equality, giving

$$\|x'_i\|^2 = \alpha^2q + 2\alpha\sqrt{q}p + p + o_n(1).$$

Similar reasoning can be used with Lemma 4.2.5 to obtain

$$\langle x'_i, x'_j \rangle = p(\alpha\sqrt{q} + 1)^2 + o_n(1).$$

Hence,

$$\lim_{n \rightarrow \infty} \langle u'_i, u'_j \rangle = \frac{p(\alpha\sqrt{q} + 1)^2}{\alpha^2q + 2\alpha\sqrt{q}p + p}.$$

If $\gamma = \frac{1}{1-p}$, we first observe

$$\frac{\sum_{j \neq i} e^{s_{ij}}}{e^\beta} = \frac{(n-1)n^{\gamma p}}{n^\gamma} = n^{\gamma(p-1)+1} - n^{\gamma(p-1)} = 1 + o_n(1),$$

thus $Z_i = (2 + o_n(1))e^\beta$. Similar calculations as in Lemma 4.2.4 yield

$$\begin{aligned} \langle u_i, \text{ATT}(u_i) \rangle &= \frac{1+p}{2} + o_n(1), \\ \langle \text{ATT}(u_i), \text{ATT}(u_i) \rangle &= \frac{1+3p}{4} + o_n(1). \end{aligned}$$

Similar calculations as in Lemma 4.2.5 give

$$\begin{aligned} \langle u_i, \text{ATT}(u_j) \rangle &= p + o_n(1), \\ \langle \text{ATT}(u_i), \text{ATT}(u_j) \rangle &= p + o_n(1). \end{aligned}$$

By substituting these results into the inner product expansions, we obtain

$$\langle u'_i, u'_j \rangle = \frac{\langle x'_i, x'_j \rangle}{\|x'_i\| \|x'_j\|} = \frac{p(\alpha\sqrt{q} + 1)^2}{\alpha^2 q + \alpha\sqrt{q}(1+p) + \frac{1+3p}{4}} + o_n(1).$$

If $\gamma > \frac{1}{1-p}$, Theorem 4.2.1 directly gives

$$\lim_{n \rightarrow \infty} \langle u_i, u_j \rangle = \langle u_i, u_j \rangle = p.$$

□

Since $p < 1$, the trailing term in the denominators of the $\gamma \geq \frac{1}{1-p}$ cases show that we have $\lim_{n \rightarrow \infty} \langle u'_i, u'_j \rangle > p$. Thus, even when scaled appropriately, attention is still inherently a contractive operator.

Remark 4.2.1: In the absence of residual connections, i.e. $\alpha = 0$, we have

$$\lim_{n \rightarrow \infty} \langle u'_i, u'_j \rangle = \begin{cases} 1 & \gamma < \frac{1}{1-p}, \\ \frac{4p}{1+3p} & \gamma = \frac{1}{1-p}, \\ p & \gamma > \frac{1}{1-p}. \end{cases}$$

The phase transition is especially clear in this case. In the subcritical regime, all token embeddings collapse onto the same direction, which implies that attention is behaving like an averaging operator, $\text{ATT}(u_i) = \frac{1}{n} \sum_j u_j$. In the supercritical regime, the geometry of the embeddings is asymptotically unchanged and attention behaves like the identity map. In some sense, this provides a reason for why residual connections are crucial in avoiding this rank collapse phenomenon.

4.3 Backpropagation

Theorem 4.3.1 (Phase Transition in Attention Gradient):

Assume there exist constants $q_1, q_2 > 0$ and $p_1, p_2 \in (0, 1)$ such that $q_1 \leq \|x_i\|^2 \leq q_2$ and $p_1 \leq \langle u_i, u_j \rangle \leq p_2$ for any $i \neq j$, with $p_1 = \langle u_i, u_j \rangle$ for some i, j . Let $\beta = \gamma \log n$ with $\gamma > 0$.

1. If $\gamma < \frac{1}{1-p_1}$,

$$\frac{1}{nd} \|\nabla_X X'\|^2 \leq \frac{4\gamma^2 (\log n)^2}{q_1 d} + o_n(1).$$

2. If $\gamma > \frac{1}{1-p_2}$,

$$\frac{1}{nd} \|\nabla_X X'\|^2 \geq \frac{1}{q_2} \left(1 - \frac{1}{d}\right) + o_n(1).$$

[Che+25]

Lemma 4.3.2: For any $i, k \in \{1, \dots, n\}$ and $u, w \in \{1, \dots, d\}$,

$$\frac{\partial(N(x_k))_w}{\partial(x_i)_u} = \delta_{ik} \frac{\delta_{wu} \|x_k\|^2 - (x_k)_w (x_k)_u}{\|x_k\|^3}$$

Proof. By the quotient rule,

$$\frac{\partial(N(x_k))_w}{\partial(x_i)_u} = \frac{\partial((x_k)_w \cdot \|x_k\|^{-1})}{\partial(x_i)_u} = \delta_{ik} \frac{\delta_{wu} \|x_k\| - (x_k)_w \cdot \frac{(x_k)_u}{\|x_k\|}}{\|x_k\|^2} = \delta_{ik} \frac{\delta_{wu} \|x_k\|^2 - (x_k)_w (x_k)_u}{\|x_k\|^3}.$$

□

Lemma 4.3.3: For any $k, j \in \{1, \dots, n\}$ and $w, v \in \{1, \dots, d\}$,

$$\begin{aligned} \frac{\partial(\text{ATT}(u_j))_v}{\partial(u_k)_w} &= \left[\left(\delta_{kj} \beta \left(\sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (u_m)_w (u_m)_v \right) + e^{\beta \langle u_j, u_k \rangle} (\beta(u_j)_w (u_k)_v + \delta_{wv}) \right) \cdot Z_j \right. \\ &\quad \left. - \left(\delta_{kj} \beta \left(\sum_{l=1}^n e^{\beta \langle u_j, u_l \rangle} (u_l)_w \right) + \beta e^{\beta \langle u_j, u_k \rangle} (u_j)_w \right) \cdot \left(\sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (u_m)_v \right) \right] \cdot Z_j^{-2} \end{aligned}$$

Proof. Recall

$$(\text{ATT}(u_j))_v = \frac{\sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (u_m)_v}{\sum_{l=1}^n e^{\beta \langle u_j, u_l \rangle}} := \frac{(y_j)_v}{Z_j}$$

The derivative of the numerator,

$$\begin{aligned} \frac{\partial(y_j)_v}{\partial(u_k)_w} &= \sum_{m=1}^n \left[\left(\frac{\partial}{\partial(u_k)_w} e^{\beta \langle u_j, u_m \rangle} \right) (u_m)_v + e^{\beta \langle u_j, u_m \rangle} \frac{\partial}{\partial(u_k)_w} (u_m)_v \right] \\ &= \sum_{m=1}^n \left[\beta \left(\delta_{kj} (u_m)_w + \delta_{km} (u_j)_w \right) e^{\beta \langle u_j, u_m \rangle} (u_m)_v + e^{\beta \langle u_j, u_m \rangle} (\delta_{km} \delta_{wv}) \right] \\ &= \left(\beta \delta_{kj} \sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (u_m)_w (u_m)_v \right) + e^{\beta \langle u_j, u_k \rangle} (\beta(u_j)_w (u_k)_v + \delta_{wv}). \end{aligned}$$

The derivative of the denominator,

$$\frac{\partial Z_j}{\partial(u_k)_w} = \sum_{l=1}^n \beta \left(\delta_{kj} (u_l)_w + \delta_{kl} (u_j)_w \right) e^{\beta \langle u_j, u_l \rangle} = \left(\delta_{kj} \beta \sum_{l=1}^n e^{\beta \langle u_j, u_l \rangle} (u_l)_w \right) + \beta e^{\beta \langle u_j, u_k \rangle} (u_j)_w.$$

Applying the quotient rule yields the desired result. □

Lemma 4.3.4: The Jacobian of the attention operator has the form

$$\left(\frac{\partial \left(\text{ATT}(N(x_j)) \right)_v}{\partial (x_i)_u} \right)_{u,v=1}^d = \|x_i\|^{-1/2} [(\mathbf{R}_1 + \mathbf{R}_2)Z_j - (\mathbf{U}_1 + \mathbf{U}_2) \otimes \mathbf{V}_j] \cdot Z_j^{-2},$$

where

$$\begin{aligned} \mathbf{R}_1 &= \delta_{ij}(\mathbf{W}_j - u_i \otimes (\mathbf{W}_j u_j)), & \mathbf{R}_2 &= e^{\beta \langle u_j, u_i \rangle} ((-u_j + \beta \mathbf{P}_{u_i} u_j) \otimes u_i + I_d), \\ \mathbf{U}_1 &= \delta_{ij} \beta (\mathbf{P}_{u_i} \mathbf{V}_j), & \mathbf{U}_2 &= \beta e^{\beta \langle u_j, u_i \rangle} (\mathbf{P}_{u_i} u_j), \\ \mathbf{V}_j &= \sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} u_m, & \mathbf{W}_j &= \sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} u_m \otimes u_m, \\ \mathbf{P}_x y &= y - \langle y, x \rangle x. \end{aligned}$$

Proof. Let

$$J_{wv} = \frac{\partial (\text{ATT}(N(x_j)))_v}{\partial (N(x_k))_w} = \frac{\partial (\text{ATT}(u_j))_v}{\partial (u_k)_w},$$

which was computed in the previous lemma. By the multivariate chain rule, we have

$$\begin{aligned} \frac{\partial \left(\text{ATT}(N(x_j)) \right)_v}{\partial (x_i)_u} &= \sum_{k=1}^n \sum_{w=1}^d J_{wv} \cdot \frac{\partial (N(x_k))_w}{\partial (x_i)_u} \\ &= \sum_{w=1}^d J_{wv} \cdot \frac{\delta_{wu} \|x_i\|^2 - (x_i)_w (x_i)_u}{\|x_i\|^3} \\ &= \|x_i\|^{-1} \left(J_{uv} - (u_i)_u \sum_{w=1}^d J_{wv} (u_i)_w \right). \end{aligned}$$

We compute

$$\begin{aligned} &\sum_{w=1}^d J_{wv} (u_i)_w \\ &= \left[\left(\delta_{ij} \beta \left(\sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} \langle u_m, u_i \rangle (u_m)_v \right) + e^{\beta \langle u_j, u_i \rangle} ((\beta \langle u_j, u_i \rangle + 1) (u_i)_v) \cdot Z_j \right. \right. \\ &\quad \left. \left. - \left(\delta_{ij} \beta \left(\sum_{l=1}^n e^{\beta \langle u_j, u_l \rangle} \langle u_l, u_i \rangle \right) + \beta e^{\beta \langle u_j, u_i \rangle} \langle u_j, u_i \rangle \right) \cdot \left(\sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (u_m)_v \right) \right] \cdot Z_j^{-2}, \end{aligned}$$

which then gives

$$\begin{aligned}
\frac{\partial \left((\text{ATT}(N(x_j)))_v \right)}{\partial (x_i)_u} &= \left[\left(\delta_{ij} \beta \left(\sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} \left((u_m)_u (u_m)_v - \langle u_m, u_i \rangle (u_m)_v (u_i)_u \right) \right) \right. \right. \\
&\quad + \left. e^{\beta \langle u_j, u_k \rangle} \left(\beta (u_j)_u (u_k)_v + \delta_{uv} - (\beta \langle u_j, u_i \rangle + 1) (u_i)_v (u_i)_u \right) \right) \cdot Z_j \\
&\quad - \left(\delta_{ij} \beta \left(\sum_{l=1}^n e^{\beta \langle u_j, u_l \rangle} \left((u_l)_u - \langle u_l, u_i \rangle (u_i)_u \right) \right) \right. \\
&\quad + \left. \beta e^{\beta \langle u_j, u_k \rangle} \left((u_j)_u - \langle u_j, u_i \rangle (u_i)_u \right) \right) \cdot \left(\sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (u_m)_v \right) \Big] \\
&\quad \cdot Z_j^{-2} \cdot \|x_i\|^{-1}.
\end{aligned}$$

When written in matrix form (i.e. the transposed Jacobian), this becomes

$$\begin{aligned}
&\left[\left[\delta_{ij} \beta (\mathbf{W}_j - u_i \otimes (W_j u_i)) + e^{\beta \langle u_j, u_k \rangle} (\beta u_j \otimes u_i + I_d - (\beta \langle u_j, u_i \rangle + 1) u_i \otimes u_i) \right] \cdot Z_j \right. \\
&\quad \left. - \left[\delta_{ij} \beta (\mathbf{V}_j - \langle \mathbf{V}_j, u_i \rangle u_i) + \beta e^{\beta \langle u_j, u_i \rangle} (u_j - \langle u_j, u_i \rangle u_i) \right] \otimes \mathbf{V}_j \right] \cdot Z_j^{-2} \cdot \|x_i\|^{-1} \\
&= \left[(\mathbf{R}_1 + e^{\beta \langle u_j, u_i \rangle} ((-u_i + \beta \mathbf{P}_{u_i} u_j) \otimes u_i + I_d)) \cdot Z_j \right. \\
&\quad \left. - \left[\delta_{ij} \beta (\mathbf{P}_{u_i} \mathbf{V}_j) + \beta e^{\beta \langle u_j, u_i \rangle} (\mathbf{P}_{u_i} u_j) \right] \otimes \mathbf{V}_j \right] \cdot Z_j^{-2} \cdot \|x_i\|^{-1} \\
&= \left[(\mathbf{R}_1 + \mathbf{R}_2) \cdot Z_j + (\mathbf{U}_1 + \mathbf{U}_2) \otimes \mathbf{V}_j \right] \cdot Z_j^{-2} \cdot \|x_i\|^{-1},
\end{aligned}$$

where the intermediate steps follow from basic properties of the tensor product. $\square$

We now prove Theorem 4.2.1.

Proof of Theorem 4.2.1. Throughout this proof, we use the fact that $\|x \otimes y\|_F = \|x\|_2 \|y\|_2$, as well as $\|A\|_F^2 = \text{tr}(A^T A)$. The subscripts will be dropped since we are exclusively using the Frobenius norm for matrices and 2-norm for vectors. First consider the case when $\gamma > \frac{1}{1-p_2}$.

1. When $i \neq j$, $\mathbf{R}_1 Z_j^{-1} = 0$. When $i = j$, we obtain (using the form of $\mathbf{R}_1$ before simplification in the previous lemma)

$$\begin{aligned}
\|\mathbf{R}_1\|Z_j^{-1} &= \left\| \beta \sum_{m=1}^n e^{\beta\langle u_j, u_m \rangle} (u_m \otimes u_m - \langle u_m, u_i \rangle u_i \otimes u_m) \right\| \cdot Z_j^{-1} \\
&= \left\| \beta \sum_{m \neq j} e^{\beta\langle u_j, u_m \rangle} (u_m \otimes u_m - \langle u_m, u_i \rangle u_i \otimes u_m) \right\| \cdot Z_j^{-1} \\
&\leq \beta \left[\sum_{m \neq j} e^{\beta p_2} (\|u_m\|^2 + \|u_m\|^2 \|u_i\|^2) \right] \cdot Z_j^{-1} \\
&= 2\beta(n-1)e^{\beta p_2} \cdot Z_j^{-1} \\
&\leq 2\gamma \log(n) \cdot n^{\gamma p_2 + 1 - \gamma} (1 + o_n(1))
\end{aligned}$$

where the last inequality follows from Lemma 4.2.2. Since $\gamma p_2 + 1 - \gamma < 0$ by assumption, this value goes to 0 as $n \rightarrow \infty$.

2. When $i = j$, we notice that $\mathbf{P}_{u_i} u_j = 0$, hence

$$\mathbf{R}_2 Z_j^{-1} = e^{\beta} Z_j^{-1} (-u_i \otimes u_i + I_d) = (I_d - u_i \otimes u_i)(1 + o_n(1))$$

When $i \neq j$, naively bounding gives

$$\begin{aligned}
\|\mathbf{R}_2\|Z_j^{-1} &\leq e^{\beta p_2} Z_j^{-1} (\|u_i\|^2 + \beta \|\mathbf{P}_{u_i} u_j\| \|u_i\| + \|I_d\|) \\
&\leq n^{\gamma(p_2-1)} (1 + 2\gamma \log n + \sqrt{d}) (1 + o_n(1)),
\end{aligned}$$

which goes to 0 as $n \rightarrow \infty$ since $\gamma > 0$ and $p_2 - 1 < 0$.

3. When $i = j$, we have that $\mathbf{U}_2 = 0$ since $\mathbf{P}_{u_i} u_i = 0$, so

$$\begin{aligned}
\|(\mathbf{U}_1 + \mathbf{U}_2) \otimes \mathbf{V}_j\| \cdot Z_j^{-2} &= Z_j^{-2} \beta \|\mathbf{P}_{u_i} \mathbf{V}_j\| \|\mathbf{V}_j\| \\
&\leq Z_j^{-2} \beta \left(\sum_{m \neq j} e^{\beta\langle u_j, u_m \rangle} \|\mathbf{P}_{u_i} u_m\| \right) \cdot \left(e^{\beta} + \sum_{m \neq j} e^{\beta\langle u_j, u_m \rangle} \|u_m\| \right) \\
&\leq 2e^{-2\beta} \beta (ne^{\beta p_2}) (e^{\beta} + ne^{\beta p_2}) (1 + o_n(1)) \\
&\leq 2\gamma \log n \cdot n^{\gamma(p_2-1)+1} \cdot (1 + n^{\gamma(p_2-1)+1}) (1 + o_n(1)),
\end{aligned}$$

which goes to 0 as $n \rightarrow \infty$ since $\gamma(p_2 - 1) + 1 < 0$. When $i \neq j$, $\mathbf{U}_1 = 0$, thus

$$\begin{aligned}
\|(\mathbf{U}_1 + \mathbf{U}_2) \otimes \mathbf{V}_j\| \cdot Z_j^{-2} &= Z_j^{-2} \beta e^{\beta\langle u_j, u_i \rangle} \|\mathbf{P}_{u_i} u_j\| \|\mathbf{V}_j\| \\
&\leq 2Z_j^{-2} \beta e^{\beta p_2} \sum_{m=1}^n e^{\beta\langle u_j, u_m \rangle} \|u_m\| \\
&\leq 2\beta e^{-2\beta} e^{\beta p_2} n e^{\beta} (1 + o_n(1)) \\
&\leq 2\gamma \log(n) n^{\gamma(p_2-1)+1} (1 + o_n(1)),
\end{aligned}$$

which goes to 0 as $n \rightarrow \infty$ since $\gamma(p_2 - 1) + 1 < 0$. Combining all cases for $\gamma > \frac{1}{1-p_2}$, we have that

$$\begin{aligned}
\left(\frac{\partial \left(\text{ATT}(N(x_j)) \right)_v}{\partial (x_i)_u} \right)_{u,v=1}^d &\leq \frac{\delta_{ij}}{\|x_i\|} (I_d - u_i \otimes u_i) (1 + o_n(1)) \\
&= \frac{\delta_{ij}}{\|x_i\|} (I_d - u_i \otimes u_i) + \mathbf{o}_n(1) + o_n(1) I_d.
\end{aligned}$$

Noticing the block-diagonal structure of $\nabla_X X'$ and that $(u_i \otimes u_i)^2 = u_i \otimes u_i$,

$$\begin{aligned}
\frac{1}{nd} \|\nabla_X X'\|^2 &= \frac{1}{nd} \left(\sum_{i=1}^n \frac{1}{\|x_i\|} \|I_d - u_i \otimes u_i\|^2 \right) + o_n(1) \\
&\geq \frac{1}{ndq_2} \left(\sum_{i=1}^n \text{tr}(I_d - u_i \otimes u_i) \right) + o_n(1) \\
&= \frac{1}{ndq_2} \left(\sum_{i=1}^n d - \text{tr}(u_i u_i^T) \right) + o_n(1) \\
&= \frac{1}{q_2} \left(1 - \frac{1}{d} \right) + o_n(1).
\end{aligned}$$

Next, consider the case when $\gamma < \frac{1}{1-p_2}$.

1. Once again, if $i \neq j$, $\mathbf{R}_1 Z_j^{-1} = 0$. When $i = j$, we first note that $\|\mathbf{P}_{u_i} u_m\| \leq 1$, since $\mathbf{P}_{u_i} u_m$ gives the component of u_m orthogonal to u_i . Thus,

$$\begin{aligned}
\mathbf{R}_1 Z_j^{-1} &= \left\| \beta \sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (u_m \otimes u_m - \langle u_m, u_i \rangle u_i \otimes u_m) \right\| \cdot Z_j^{-1} \\
&= \beta \left\| \sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} (\mathbf{P}_{u_i} u_m) \otimes u_m \right\| \cdot Z_j^{-1} \\
&\leq \beta Z_j^{-1} \sum_{m=1}^n e^{\beta \langle u_j, u_m \rangle} \|\mathbf{P}_{u_i} u_m\| \|u_m\| \\
&\leq \beta
\end{aligned}$$

2. As computed earlier, $\|I_d - u_i \otimes u_i\| = \sqrt{d-1}$, giving

$$\begin{aligned}
\|\mathbf{R}_2\| Z_j^{-1} &\leq Z_j^{-1} e^{\beta \langle u_j, u_i \rangle} \left\| (-u_i + \beta \mathbf{P}_{u_i} u_j) \otimes u_i + I_d \right\| \\
&\leq Z_j^{-1} e^{\beta \langle u_j, u_i \rangle} (\|I_d - u_i \otimes u_i\| + \beta \|\mathbf{P}_{u_i} u_j\| \|u_i\|) \\
&\leq Z_j^{-1} e^{\beta \langle u_j, u_i \rangle} (\sqrt{d-1} + \beta).
\end{aligned}$$

3. Note that $\|\mathbf{V}_j\| \leq Z_j$, and $\|\mathbf{P}_{u_i} \mathbf{V}_j\| \leq \|\mathbf{V}_j\| = Z_j$.

$$\begin{aligned}
\|(\mathbf{U}_1 + \mathbf{U}_2) \otimes \mathbf{V}_j\| \cdot Z_j^{-2} &\leq (\|\mathbf{U}_1\| + \|\mathbf{U}_2\|) \cdot Z_j^{-1} \\
&\leq Z_j^{-1} \delta_{ij} \beta \|\mathbf{P}_{u_i} \mathbf{V}_j\| + \beta e^{\beta \langle u_j, u_i \rangle} \|\mathbf{P}_{u_i} u_j\| \cdot Z_j^{-1} \\
&\leq \beta (\delta_{ij} + e^{\beta \langle u_j, u_i \rangle} Z_j^{-1}).
\end{aligned}$$

All together, this gives

$$\begin{aligned}
&\left\| \left(\frac{\partial \left((\text{ATT}(N(x_j)))_v \right)}{\partial (x_i)_u} \right)_{u,v=1}^d \right\| \\
&\leq \frac{1}{\|x_i\|} (\delta_{ij} \beta + Z_j^{-1} e^{\beta \langle u_j, u_i \rangle} (\sqrt{d-1} + \beta) + \beta (\delta_{ij} + e^{\beta \langle u_j, u_i \rangle} Z_j^{-1})) \\
&= \frac{1}{\|x_i\|} (2\beta \delta_{ij} + (2\beta + \sqrt{d}) e^{\beta \langle u_j, u_i \rangle} Z_j^{-1}).
\end{aligned}$$

Using the properties of Z_j as shown in Lemma 4.2.2,

$$\begin{aligned}
&\frac{1}{nd} \|\nabla_X X'\|^2 \\
&\leq \frac{1}{nd} \sum_{i,j=1}^n \frac{1}{\|x_i\|} \left(4\beta^2 \delta_{ij} + 4\beta \delta_{ij} (2\beta + \sqrt{d}) \frac{e^{\beta \langle u_j, u_i \rangle}}{Z_j} + (2\beta + \sqrt{d})^2 \cdot \frac{e^{2\beta \langle u_j, u_i \rangle}}{Z_j^2} \right) \\
&\leq \frac{1}{ndq_1} \left(n \cdot 4\beta^2 + O((\log n)^2) \left(\sum_{j=1}^n \frac{e^\beta}{Z_j} \right) + O((\log n)^2) \left(\sum_{i,j=1}^n \frac{e^{2\beta \langle u_j, u_i \rangle}}{Z_j^2} \right) \right) \\
&\leq \frac{(1 + o_n(1))}{ndq_1} \left(4n\beta^2 + O((\log n)^2) \frac{n^{\gamma+1}}{n^{\gamma p_1+1}} + O((\log n)^2) \cdot \sum_{j=1}^n \left(\frac{n^\gamma}{n^{\gamma p_1+1}} \cdot \sum_{i=1}^n \frac{e^{\beta \langle u_j, u_i \rangle}}{Z_j} \right) \right) \\
&= \frac{(1 + o_n(1))}{dq_1} \left(4\beta^2 + O((\log n)^2) n^{\gamma(1-p_1)-1} + O((\log n)^2) n^{\gamma(1-p_1)-1} \right) \\
&= 4 \frac{\gamma^2 (\log n)^2}{dq_1} + o_n(1),
\end{aligned}$$

since $\gamma(1 - p_1) - 1 < 0$. □

5 ATTENTION AT INITIALIZATION

A natural next step after studying the equicorrelated setting is to consider the fully random case, where the inputs no longer share a common correlation structure and instead behave like independent high-dimensional noise. This serves as a baseline model in which any structure in the attention scores comes purely from random fluctuations, most importantly in the case of a newly initialized, pre-trained model. The purpose of this regime is not to improve upon the equicorrelated analysis, but to provide a contrasting baseline.

5.1 Independent Entries

To identify a potential scaling for the full model, we first restrict ourselves to the simpler independent case. Consider a score vector $s = (s_1, \dots, s_n)$ with independent and identically distributed entries $s_i \sim \mathcal{N}(0, 1)$. If s_m is the maximum score, the maximum attention weight is given by

$$a_m = \frac{e^{\beta s_m}}{\sum_{k=1}^n e^{\beta s_k}} = \left(1 + \frac{\sum_{k \neq m} e^{\beta s_k}}{e^{\beta s_m}} \right)^{-1}$$

For large n , we expect

$$\sum_{k \neq m} e^{\beta s_k} \sim n \mathbb{E}[e^{\beta s_1}] = n e^{\beta^2/2} = \exp(\log n + \beta^2/2).$$

It can also be shown that $s_m \sim \sqrt{2 \log n}$ [LLR83] and hence $e^{\beta s_m} \sim e^{\beta \sqrt{2 \log n}}$. A scaling for which $a_m \in (0, 1)$ must then have

$$\log n + \frac{\beta^2}{2} = \beta \sqrt{2 \log n},$$

which yields $\beta_n = \sqrt{2 \log n}$ after solving. Although a rough asymptotic argument, this identifies $\beta_n \asymp \sqrt{\log n}$ as the critical scaling factor under this i.i.d. assumption.

5.2 Correlated Entries

We now generalize to a model that resembles the attention mechanism more closely. Consider the $n \times n$ score matrix obtained from the multiplication

$$S = \left[\frac{(XW_Q)(XW_K)^T}{\sqrt{d}} \right],$$

where $X \in \mathbb{R}^{n \times d}$ is a random matrix with i.i.d. $\mathcal{N}(0, 1)$ components, and $W_Q, W_K \in \mathbb{R}^{d \times d}$ are random matrices with i.i.d. $\mathcal{N}(0, 1/d)$ entries.

In this context, S is a random object representing what score values may be produced by a randomly initialized model prior to training. For this reason, the score scaling β may equivalently be viewed as a scaling factor for the variance of i.i.d. $\mathcal{N}(0, \beta/d)$ entries in the W_Q and W_K matrices. This analysis is therefore also particularly useful in the effort of counteracting pathological behaviour at the beginning of training, which could in the worst case lead to an untrainable model. Since the methods used in this section are closely related to those of statistical physics, we will rely on asymptotic approximations and high-probability concentration results, as opposed to the exact results found in other sections.

Definition 5.2.1: Given attention weights $a_i \in \mathbb{R}^n$, define the *inverse participation ratio*

$$Y_i^{(2)} = \sum_{j=1}^n a_{ij}^2.$$

The inverse participation ratio heuristically satisfies

$$Y_i^{(2)} = \frac{1}{\text{effective support size of } a_i},$$

since $Y_i^{(2)} = 1$ if a_i has a single coordinate with value 1, and $Y_i^{(2)} = \frac{1}{n}$ if a_i is uniform. We are specifically interested in a scaling β that avoids these two cases in the large- n limit.

Theorem 5.2.2 (Phase Transition for IPR):

Under the scaling $\beta_n = \gamma\sqrt{\log n}$ where $\gamma > 0$ is a constant, the inverse participation ratio approximately satisfies

$$\lim_{n \rightarrow \infty} \mathbb{E}[Y_i^{(2)}(\beta)] = \begin{cases} 0, & \gamma < \sqrt{2}, \\ 1 - \frac{\sqrt{2}}{\gamma}, & \gamma > \sqrt{2} \end{cases}$$

given that the embedding dimension d is sufficiently large. Thus, under this model, the critical scaling is of order $\beta \asymp \sqrt{\log n}$ with phase transition at $\beta = \sqrt{2 \log n}$. [GG25]

Lemma 5.2.3 (Concentration of Token Similarity): For two rows x_i, x_j of the embedding matrix X , define the *token similarity* $q_{ij} = \langle x_i, x_j \rangle / d$. Then,

$$q_{ij} \xrightarrow{p} \begin{cases} 1, & i = j, \\ 0, & i \neq j. \end{cases}$$

In particular, we may assume approximate deterministic equality for large d , such as those used in modern language models.

Proof. For $i = j$, we have

$$q_{ii} = \frac{1}{d} \sum_{k=1}^d (x_i)_k^2.$$

By the Law of Large Numbers, $q_{ii} \xrightarrow{p} \mathbb{E}[(x_i)_1^2] = 1$, since $(x_i)_k \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1)$. For $i \neq j$,

$$\mathbb{E}[q_{ij}] = \frac{1}{d} \sum_{k=1}^d \mathbb{E}[(x_i)_k (x_j)_k] = 0,$$

and

$$\text{Var}(q_{ij}) = \frac{1}{d} \cdot \mathbb{E}[(x_i)_k^2 (x_j)_k^2] = \frac{1}{d}.$$

By Chebyshev's Inequality, $q_{ij} \rightarrow 0$. Moreover, sharper concentration results (see [Ver26]) imply that these quantities are tightly concentrated with high probability for large d . $\square$

Lemma 5.2.4 (Stein's Lemma): Let $Z = (Z_1, \dots, Z_n)^T \sim \mathcal{N}(0, V)$ and $g : \mathbb{R}^n \rightarrow \mathbb{R}$ such that $g \in C^1$ and $\mathbb{E}[\|g(Z)\|], \mathbb{E}[\|\nabla g(Z)\|] < \infty$. Then,

$$\mathbb{E}[Z_i g(Z)] = \sum_{j=1}^n \text{Cov}(Z_i, Z_j) \cdot \mathbb{E}\left[\frac{\partial}{\partial Z_j} g(Z)\right]$$

Proof. The pdf of Z has the form $p(z) = C \exp[-\frac{1}{2}z^T V^{-1}z]$. Thus,

$$\begin{aligned} \nabla p(z) &= -V^{-1}z p(z) \\ \Rightarrow z p(z) &= -V \nabla p(z) \\ \Rightarrow z_i p(z) &= -\sum_{j=1}^n V_{ij} \frac{\partial}{\partial z_j} p(z). \end{aligned}$$

Substituting the above equation into the calculation of expectation,

$$\mathbb{E}[Z_i g(Z)] = \int_{\mathbb{R}^n} z_i g(z) p(z) dz = -\sum_{j=1}^n V_{ij} \int_{\mathbb{R}^n} g(z) \frac{\partial}{\partial z_j} p(z) dz.$$

We then apply integration by parts along the dimension of z_j , noting that $g(z)$ vanishes when $z_j \rightarrow \pm\infty$ by assumption of finite expectation. This gives

$$\mathbb{E}[Z_i g(Z)] = \sum_{j=1}^n V_{ij} \int_{\mathbb{R}^n} \frac{\partial}{\partial z_j} g(z) p(z) dz = \sum_{j=1}^n \text{Cov}(Z_i, Z_j) \cdot \mathbb{E}\left[\frac{\partial}{\partial z_j} g(Z)\right].$$

□

Lemma 5.2.5 (Derivative of Scaled Softmax): Let $S \in \mathbb{R}^{n \times n}$. For any $i, j, k \in \{1, \dots, n\}$,

$$\frac{\partial a_{ij}}{\partial s_{ik}} = \frac{\partial}{\partial s_{ik}} \left(\frac{e^{h\beta s_{ij}}}{\sum_{l=1}^n e^{h\beta s_{il}}} \right) = h\beta a_{ij} (\delta_{jk} - a_{ik}).$$

Proof. The derivative of the numerator is

$$\frac{\partial}{\partial s_{ik}} e^{h\beta s_{ij}} = h\beta e^{h\beta s_{ij}} \delta_{jk}.$$

The derivative of the denominator Z_i is

$$\frac{\partial}{\partial s_{ik}} Z_i = h\beta e^{h\beta s_{ik}}.$$

By quotient rule,

$$\frac{\partial}{\partial s_{ik}} \left(\frac{e^{h\beta s_{ij}}}{\sum_{l=1}^n e^{h\beta s_{il}}} \right) = \frac{h\beta e^{h\beta s_{ij}} \delta_{jk} Z - e^{h\beta s_{ij}} h\beta e^{h\beta s_{ik}}}{Z^2} = h\beta a_{ij} (\delta_{jk} - a_{ik}).$$

□

Lemma 5.2.6: For large d , we have

$$s_{ij} \sim \mathcal{N}(0, 1),$$

with covariance structure

$$\text{Cov}(s_{ij}, s_{pq}) = q_{ip}q_{jq}.$$

Proof. By definition,

$$\begin{aligned} s_{ij} &= \frac{1}{\sqrt{d}} \sum_{k=1}^d (W_Q x_i)_k (W_K x_j)_k \\ &= \frac{1}{\sqrt{d}} \sum_{k=1}^d \sum_{a,b=1}^d (W_Q)_{ka} (x_i)_a (W_K)_{kb} (x_j)_b \\ &:= \frac{1}{\sqrt{d}} \sum_{k=1}^d D_k. \end{aligned}$$

First, $\mathbb{E}[D_k] = 0$ since all entries of W_K and W_Q have expectation 0. As for variance,

$$\begin{aligned} \mathbb{E}[D_k^2] &= \mathbb{E} \left[\left(\sum_{a=1}^d (W_Q)_{ka} (x_i)_a \right)^2 \right] \cdot \mathbb{E} \left[\left(\sum_{b=1}^d (W_K)_{kb} (x_j)_b \right)^2 \right] \\ &= \left(\sum_{a=1}^d \mathbb{E}[(W_Q)_{ka}^2] (x_i)_a^2 \right) \left(\sum_{b=1}^d \mathbb{E}[(W_K)_{kb}^2] (x_j)_b^2 \right) \\ &= \frac{1}{d^2} \|x_i\|^2 \|x_j\|^2 \end{aligned}$$

We have that $\frac{\|x_k\|^2}{d} \xrightarrow{p} 1$ by Lemma 5.2.3, hence by the Central Limit Theorem, $s_{ij} \xrightarrow{d} \mathcal{N}(0, 1)$. A similar calculation for covariance yields

$$\begin{aligned} \text{Cov}(s_{ij}, s_{pq}) &= \mathbb{E}[s_{ij} s_{pq}] \\ &= \frac{1}{d} \sum_{k,a,b,l,\alpha,\beta=1}^d (x_i)_a (x_j)_b (x_p)_\alpha (x_q)_\beta \mathbb{E}[(W_Q)_{ka} (W_Q)_{l\alpha}] \mathbb{E}[(W_K)_{kb} (W_K)_{l\beta}] \\ &= \frac{1}{d} \sum_{k,a,b=1}^d (x_i)_a (x_j)_b (x_p)_a (x_q)_b \mathbb{E}[(W_Q)_{ka}^2] \mathbb{E}[(W_K)_{ka}^2] \\ &= \frac{1}{d^3} \sum_{k=1}^d \langle x_i, x_p \rangle \langle x_j, x_q \rangle \\ &= q_{ip} q_{jq}. \end{aligned}$$

□

Lemma 5.2.7: Define the function

$$\Phi_n(\beta, h) = \mathbb{E}[\log Z_i(\beta, h)], \quad Z_i(\beta, h) = \sum_{j=1}^n e^{h\beta s_{ij}}.$$

Then,

$$\lim_{n \rightarrow \infty} \mathbb{E}[Y_i^{(2)}] = 1 - \lim_{h \rightarrow 1} \lim_{n \rightarrow \infty} \frac{\partial}{\partial h} \frac{1}{\beta^2} \Phi_n(\beta, h).$$

Proof. We first note that we can safely swap derivative and expectation in the case of

$$\frac{\partial}{\partial h} \Phi_n(\beta, h) = \mathbb{E} \left[\frac{\partial}{\partial h} \log Z_i(\beta, h) \right] = \mathbb{E} \left[\frac{\sum_{j=1}^n \beta s_{ij} e^{h\beta s_{ij}}}{\sum_{k=1}^n e^{h\beta s_{ik}}} \right] = \beta \sum_{j=1}^n \mathbb{E}[s_{ij} a_{ij}],$$

where we temporarily use the notation $a_{ij} = \exp(h\beta s_{ij})/Z_i(\beta, h)$. Since this is a sum of elements of the form $\mathbb{E}[s_{ij}g(s_i)]$ where $s_i \in \mathbb{R}^n$ is normally distributed, we apply Stein's Lemma to obtain

$$\begin{aligned} \frac{\partial}{\partial h} \Phi_n(\beta, h) &= \beta \sum_{j=1}^n \sum_{k=1}^n \text{Cov}(s_{ij}, s_{ik}) \mathbb{E} \left[\frac{\partial a_{ij}}{\partial s_{ik}} \right] \\ &= \beta^2 \sum_{j=1}^n \sum_{k=1}^n q_{ii} q_{jk} \mathbb{E}[h\beta a_{ij}(\delta_{jk} - a_{ik})]. \end{aligned}$$

Define $Y_{i,h}^{(2)} = \sum_{j=1}^n a_{ij}^2$. By Lemma 5.2.3, $q_{ii} \xrightarrow{p} 1$ and $q_{ij} \xrightarrow{p} 0$ for $i \neq j$. This gives

$$\begin{aligned} \frac{\partial}{\partial h} \Phi_n(\beta, h) &\approx \beta^2 h \sum_{j=1}^n \sum_{k=1}^n \delta_{jk} \mathbb{E}[a_{ij} \delta_{jk} - a_{ij} a_{ik}] \\ &= \beta^2 h \sum_{j=1}^n \mathbb{E}[a_{ij} - a_{ij}^2] \\ &= \beta^2 h \left(1 - \mathbb{E}[Y_{i,h}^{(2)}] \right). \end{aligned}$$

Rearranging and taking limits gives

$$\lim_{n \rightarrow \infty} \mathbb{E}[Y_i^{(2)}] = 1 - \lim_{h \rightarrow 1} \lim_{n \rightarrow \infty} \frac{\partial}{\partial h} \frac{1}{\beta^2} \Phi_n(\beta, h).$$

□

We are now left with computing $\Phi_n(\beta, h) = \mathbb{E}[\log Z_i(\beta, h)]$. This is a problem similar to the well-studied Random Energy Model from statistical physics, for which methods exist. In particular, we use the replica method, which is based on the applicaiton of the limit

$$\mathbb{E}[\log Z] = \lim_{t \rightarrow 0^+} \frac{\mathbb{E}[Z^t] - 1}{t}.$$

In general, the replica method is not fully mathematically rigorous, especially in its use of analytic continuation and assumptions about the structure of overlap matrices. Despite this, it is able to provide accurate predictions in a wide range of disordered systems and high-dimensional probabilistic models.

Proof of Theorem 5.2.2. First, we consider the replicated system

$$\mathbb{E}[Z_i^t(\beta, h)] = \mathbb{E} \left[\left(\sum_{k=1}^n e^{h\beta s_{ik}} \right)^t \right] = \sum_{k_1, \dots, k_t=1}^n \mathbb{E} \left[\exp \left(h\beta \sum_{a=1}^t s_{ik_a} \right) \right].$$

Using the moment-generating function of the multivariate normal distribution,

$$\begin{aligned} \mathbb{E}[Z_i^t(\beta, h)] &= \sum_{k_1, \dots, k_t=1}^n \exp \left(\frac{h^2 \beta^2}{2} \sum_{a,b=1}^t \mathbb{E}[s_{ik_a} s_{ik_b}] \right) \\ &= \sum_{k_1, \dots, k_t=1}^n \exp \left(\frac{h^2 \beta^2}{2} \sum_{a,b=1}^t \delta_{k_a k_b} \right) \end{aligned}$$

Define the empirical overlap matrix Q with entries $Q_{ab} = \delta_{k_a k_b}$. Among all n^t possible configurations of replica states, there are

$$S(Q) = \sum_{k_1, \dots, k_t=1}^n \prod_{a,b=1}^t \mathbb{1}(Q_{ab} = \delta_{k_a k_b})$$

with a specific overlap matrix Q . Thus,

$$\mathbb{E}[Z_i^t(\beta, h)] = \sum_Q S(Q) \exp \left(\frac{h^2 \beta^2}{2} \sum_{a,b=1}^t Q_{ab} \right) = \sum_Q \exp \left(\log S(Q) + \frac{h^2 \beta^2}{2} \sum_{a,b=1}^t Q_{ab} \right).$$

For large n , this sum is dominated by overlap structure Q that maximizes the exponent, in accordance with the Laplace principle. To evaluate this dominant contribution, we restrict attention to a class of structured overlap matrices, which is the *1-RSB Ansatz* step in the Replica method.

By a symmetric argument, we assume that for this maximizing Q , the t states across the replicas are divided into $\frac{t}{x}$ groups of x elements each, where replicas within the same group occupy the same state, i.e. $k_a = k_b$ if and only if replicas a and b are in the same group.

Although x will ultimately be extended to real values via analytic continuation, we temporarily assume that relevant quantities are integers in order to analyze the system combinatorically. For any maximizing Q under this assumption,

$$S(Q) = n \cdot (n-1) \cdots \left(n - \frac{t}{x} + 1 \right) \cdot \frac{t!}{(x!)^{t/x} (t/x!)},$$

so $\log S(Q) \approx \frac{t}{x} \log n$ for large n , since $t, x \ll n$. Also,

$$\sum_{a,b=1}^t Q_{ab} = \sum_{\text{groups}} x^2 = tx,$$

which gives after Taylor expansion

$$\frac{\mathbb{E}[Z_i^t(\beta, h)] - 1}{t} = \frac{1}{x} \log n + \frac{h^2 \beta^2}{2} x + O(t^2).$$

Determining the value of x corresponding to this maximizing Q reduces, in this representation, to a minimization over x , the reason for which lies beyond the scope of this derivation. In the analytic continuation with $t \rightarrow 0^+$, we minimize over $x \in (0, 1)$, giving

$$x = \min \left\{ \frac{1}{h\beta} \sqrt{2 \log n}, 1 \right\}.$$

Using this,

$$\Phi_n(\beta, h) = \lim_{t \rightarrow 0^+} \left[\frac{\mathbb{E}[Z_i^t(\beta, h)] - 1}{t} \right] = \begin{cases} \log n + \frac{h^2 \beta^2}{2}, & \beta < \frac{\sqrt{2 \log n}}{h}, \\ h\beta \sqrt{2 \log n}, & \beta > \frac{\sqrt{2 \log n}}{h}. \end{cases}$$

We now assume $\beta_n = \gamma \sqrt{\log n}$. Thus,

$$\frac{\partial}{\partial h} \frac{1}{\beta^2} \Phi_n(\beta, h) = \begin{cases} h, & \gamma < \sqrt{2} \\ \frac{h\sqrt{2}}{\gamma}, & \gamma > \sqrt{2}. \end{cases}$$

By the previous lemma,

$$\lim_{n \rightarrow \infty} \mathbb{E}[Y_i^{(2)}(\beta)] = \begin{cases} 0, & \gamma < \sqrt{2}, \\ 1 - \frac{\sqrt{2}}{\gamma}, & \gamma > \sqrt{2}. \end{cases}$$

□

Corollary 5.2.1 (Update Average Cosine Similarity): Define the *updated token matrix*

$$Y = AXW_V,$$

where $A = \text{softmax}(S)$, applied row-wise, and $W_V \in \mathbb{R}^{d \times d}$ is a random matrix with i.i.d. $\mathcal{N}(0, 1/d)$ entries. For sufficiently large embedding dimension d , the updated token similarity matrix approximately satisfies

$$\lim_{n \rightarrow \infty} \mathbb{E} \left[\frac{YY^T}{d} \right] = \begin{cases} 0, & \gamma < \sqrt{2}, \\ \left(1 - \frac{\sqrt{2}}{\gamma}\right) I_d, & \gamma > \sqrt{2}. \end{cases}$$

In other words, $\|y_i\|^2/d$

...

Proof. First observe that $\mathbb{E}[W_V W_V^T] = I_d$ by a simple calculation and for large d , $XX^T \approx I_n$ by Lemma 5.2.3. Since W_V and a_i are independent,

$$\begin{aligned} \mathbb{E}\left[\frac{\langle y_i, y_j \rangle}{d}\right] &= \frac{1}{d} \cdot \mathbb{E}\left[(a_i X W_V)(a_j X W_V)^T\right] \\ &= \frac{1}{d} \cdot \mathbb{E}\left[a_i X \mathbb{E}[W_V W_V^T] X^T a_j^T\right] \\ &\approx \mathbb{E}[a_i a_j^T] \\ &= \delta_{ij} Y_i^2 \end{aligned}$$

□

6 DETERMINISTIC SCORE THEORY

This section contains results found in [HK26].

6.1 Scaling Order

We now concern ourselves with a fixed score vector $s \in \mathbb{R}^n$. The following analysis is independent of the attention mechanism, in the sense that the scores are taken as given, where the specific token vectors and weight matrices used to obtain these scores are not of concern. Let s^* be the maximum of such scores, and define the *gaps* $\Delta_j = s^* - s_j$ from each score to the maximum. We refer to scores as *competitors*, as they are “competing” for the attention weight that is to be distributed by softmax. Thus, the attention weight of any given competitor can be rewritten as

$$a_j(\beta) = \frac{e^{-\beta \Delta_j}}{\sum_{k=1}^n e^{-\beta \Delta_k}},$$

which is equivalent to scaled softmax applied to the new score vector Δ with maximum 0. We define $Z(\beta) = \sum_{j=1}^n \exp(-\beta \Delta_j)$ as in previous sections, while noting that it is defined on gap scores and not raw scores in this context.

Definition 6.1.1 (CGF): The *cumulative gap-counting function* (CGF) is the number of competitors with gap at most t from the maximum,

$$N(t) = \sum_{j=1}^n \mathbb{1}(\Delta_j \leq t).$$

Now, $Z(\beta)$ can be rewritten as the Laplace transform of the CGF, since

$$Z(\beta) = \sum_{j=1}^n e^{-\beta \Delta_j} = \sum_{j=1}^n \beta \int_0^\infty \mathbb{1}(\Delta_j \leq t) e^{-\beta t} dt = \beta \int_0^\infty e^{-\beta t} N(t) dt.$$

This is clearly well-defined on finite n since $N(t) \leq n$. However, we immediately observe that if $N(t)$ is not of exponential order, the attention weights tend toward uniformity.

Definition 6.1.2: The *upper tail accumulation scale* Λ is the smallest exponent for which $N(t) \leq \exp(\Lambda t)$ for all $t > 0$. Equivalently,

$$\Lambda = \sup_{t>0} \frac{\log N(t)}{t}.$$

When $\Lambda < \infty$, a gap Δ_j is called a *contact gap* if $N(\Delta_j) = e^{\Lambda \Delta_j}$. Let Δ_Λ be the largest of all such contact gaps and define the *contact accumulation exponent*

$$\alpha = \frac{\log N(\Delta_\Lambda)}{\log n} \in [0, 1],$$

which satisfies $N(\Delta_\Lambda) = n^\alpha$.

Lemma 6.1.3: The Shannon entropy $H(\beta) = -\sum_{j=1}^n a_j(\beta) \log(a_j(\beta))$ satisfies

$$H(\beta) = \log Z(\beta) - \beta \frac{Z'(\beta)}{Z(\beta)}$$

Proof. Recall that $\sum_{j=1}^n a_j(\beta) = 1$. Thus,

$$H(\beta) = -\sum_{j=1}^n a_j(\beta)(-\beta \Delta_j - \log Z(\beta)) = \log Z(\beta) + \sum_{j=1}^n \beta \Delta_j \frac{e^{-\beta \Delta_j}}{Z(\beta)} = \log Z(\beta) - \beta \frac{Z'(\beta)}{Z(\beta)}.$$

□

Theorem 6.1.4 (Sequence Collapse Criterion):

For each $n \geq 2$, let $s^{(n)} \in \mathbb{R}^n$ be a score vector of length n with corresponding upper tail accumulation scale Λ_n . For any positive sequence (β_n) ,

1. If $\beta_n/\Lambda_n \rightarrow 0$, then *top-two collapse* holds, i.e.

$$\lim_{n \rightarrow \infty} D_n(\beta_n) = 0,$$

where $D_n(\beta_n)$ is the difference between the largest and second largest attention weights $a^{(n)}(\beta_n)$.

2. If $\beta_n/\Lambda_n \rightarrow \infty$, then *entropy collapse* holds, i.e.

$$\lim_{n \rightarrow \infty} H_n(\beta_n) = 0,$$

where is the Shannon entropy of the attention weights $a^{(n)}(\beta_n)$.

[HK26]

Proof. Let $r_n = \beta_n/\Lambda_n$.

1. Let $\Delta_{\min}^{(n)}$ be the smallest non-zero gap, which necessarily has $\Delta_{\min}^{(n)} \leq \Delta_\Lambda^{(n)}$. The corresponding difference in attention weights can be bounded using $1 - e^{-x} \leq x$, giving

$$D_n(\beta_n) = \frac{1 - \exp(-\beta_n \Delta_{\min}^{(n)})}{Z(\beta_n)} \leq \frac{\beta_n \Delta_{\min}^{(n)}}{Z(\beta_n)} \leq \frac{\beta_n \Delta_{\Lambda}^{(n)}}{Z(\beta_n)}.$$

By definition, we have that $\log\left(N\left(\Delta_{\Lambda}^{(n)}\right)\right) = \Lambda_n \Delta_{\Lambda}^{(n)} = \beta_n \Delta_{\Lambda}^{(n)} / r_n$. Substituting into and rearranging the definition of α_n gives $\beta_n \Delta_{\Lambda}^{(n)} = r_n \alpha_n \log n$. Since $e^{-\beta_n t}$ is monotone on the $N\left(\Delta_{\Lambda}^{(n)}\right)$ gaps with $\Delta_j \leq \Delta_{\Lambda}^{(n)}$, we have that

$$\begin{aligned} Z(\beta_n) &\geq N\left(\Delta_{\Lambda}^{(n)}\right) \exp\left(-\beta_n \Delta_{\Lambda}^{(n)}\right) \\ &= \exp\left(\beta_n \Delta_{\Lambda}^{(n)} / r_n\right) n^{r_n \alpha_n} \\ &= n^{\alpha_n(1-r_n)}. \end{aligned}$$

Thus,

$$D_n(\beta_n) \leq \frac{\beta_n \Delta_{\Lambda}^{(n)}}{n^{\alpha_n(1-r_n)}} = r_n \alpha_n \log n \cdot n^{-\alpha_n(1-r_n)}.$$

Since $te^{-t} \leq e^{-1}$ for all $t \in \mathbb{R}$, setting $t = (1 - r_n) \cdot \alpha_n \log n$ yields

$$D_n(\beta_n) \leq \frac{r_n}{1 - r_n} \cdot e^{-1},$$

which goes to 0 as $n \rightarrow \infty$ since $r_n \rightarrow 0$ by assumption.

2. First observe that $Z(\beta_n) \geq 1$, since the gap between the maximum score and itself is always 0. Using the previously established integral form for $Z(\beta_n)$ as well as the fact that $N(t) \leq \exp(\Lambda^{(n)} t)$ by definition,

$$Z(\beta_n) = \beta_n \int_0^\infty e^{-\beta_n t} N(t) dt \leq \beta_n \int_0^\infty e^{-(\beta_n - \Lambda_n)t} dt = \frac{\beta_n}{\beta_n - \Lambda_n} = \frac{r_n}{r_n - 1},$$

which goes to 1 as $n \rightarrow \infty$ since $r_n \rightarrow \infty$. Thus, $Z(\beta_n) \rightarrow 1$.

For any n , differentiating the integral form with respect to β gives

$$Z'(\beta) = \int_0^\infty e^{-\beta t} N(t) dt + \beta \frac{d}{d\beta} \left(\int_0^\infty e^{-\beta t} N(t) dt \right) = \beta^{-1} Z(\beta) - \beta \int_0^\infty te^{-\beta t} N(t) dt,$$

which implies that

$$\begin{aligned} -\beta Z'(\beta) &= \beta^2 \int_0^\infty te^{-\beta t} N(t) dt - Z(\beta) \\ &\leq \beta^2 \int_0^\infty te^{-(\beta - \Lambda_n)t} dt - Z(\beta) \\ &= \left(\frac{\beta}{\beta - \Lambda_n} \right)^2 - \left(\frac{\beta}{\beta - \Lambda_n} \right). \end{aligned}$$

Since $(\beta_n/(\beta_n - \Lambda_n)) \rightarrow 1$ and $-\beta Z'(\beta) = \beta \sum_{j=1}^n \Delta_j \exp(-\beta \Delta_j) \geq 0$, we have that $-\beta Z'(\beta) \rightarrow 0$. By Lemma 6.1.3,

$$\lim_{n \rightarrow \infty} H(\beta_n) = \lim_{n \rightarrow \infty} \left[\log Z(\beta_n) + \frac{1}{Z(\beta_n)} (-\beta(Z'(\beta))) \right] = 0.$$

□

This establishes Λ_n as the order of the critical scaling for $s^{(n)}$. Denote $a_n \asymp b_n$ if there exists constants $c, C > 0$ such that for large enough n , $cb_n \leq a_n \leq Cb_n$.

Corollary 6.1.1 (Critical Scaling Exponent): Let $\xi_\Lambda > 0$. If $\Lambda_n \asymp (\log n)^{\xi_\Lambda}$, then ξ_Λ is unique and any non-collapsing scaling β_n must have $\beta_n \asymp (\log n)^{\xi_\Lambda}$.

Furthermore, if $\alpha_n \asymp (\log n)^{\xi_\alpha}$ and $\Delta_\Lambda^{(n)} \asymp (\log n)^{\xi_\Delta}$, then

$$\xi_\Lambda = \xi_\alpha - \xi_\Delta + 1.$$

Proof. If $\Lambda_n \asymp (\log n)^{\xi'}$ for $\xi' > 0$, then $1 \asymp (\log n)^{\xi_\Lambda - \xi'}$, which is only possible when $\xi_\Lambda = \xi'$. For a non-collapsing scaling, we must have $\beta_n/\Lambda_n \asymp 1$, which is only possible if $\beta_n \asymp (\log n)^{\xi_\Lambda}$.

Since $\Lambda_n = (\alpha_n \log n)/\Delta_\Lambda^{(n)}$,

$$\Lambda_n \asymp \frac{(\log n)^{\xi_\alpha} (\log n)}{(\log n)^{\xi_\Delta}} = (\log n)^{\xi_\alpha - \xi_\Delta + 1},$$

and by uniqueness of this exponent must mean that $\xi_\Lambda = \xi_\alpha - \xi_\Delta + 1$. □

Consider the Simplex case from Section 4. For any row i of the $n \times n$ similarity matrix,

$$a_{ij} = \begin{cases} q, & i = j, \\ p, & i \neq j. \end{cases}$$

The gap counting function is given by

$$N(t) = \begin{cases} 1, & t < q - p, \\ n, & t \geq q - p \end{cases}$$

which immediately gives $\Lambda_n = (q - p)^{-1} \log(n)$, $\alpha_n = 1$, $\Delta_\Lambda^{(n)} = q - p$. Thus, $\xi_\Lambda = 1$, $\xi_\alpha = 0$, and $\xi_\Delta = 0$, which satisfy the relation derived in Corollary 6.1.1. Furthermore, we see that the non-collapsing scaling has $\beta_n \asymp \log n$, which validates Theorem 4.2.1.

6.2 Rank-Collapse Boundary

Proposition 6.2.1 (Ties at the maximum): Let $s \in \mathbb{R}^n$. If $N(0) \geq 2$, then $\Lambda = \infty$ and for every $\beta > 0$,

$$\|a(\beta)\|_\infty \leq \frac{1}{N(0)}, \quad H(\beta) \geq \log N(0).$$

In particular, for a sequence of score vectors $(s^{(n)})$ with $s^{(n)} \in \mathbb{R}^n$ and $N(0) \geq 2$ for infinitely many n , no sequence X_n exists such that $\frac{\beta_n}{X_n} \rightarrow \infty$ implies that $H(\beta_n) \rightarrow 0$.

Proof. Since $N(0) \geq 2$ and N is non-decreasing, we have that

$$\Lambda = \sup_{t>0} \frac{\log N(t)}{t} \geq \sup_{t>0} \frac{\log 2}{t} = \infty.$$

The top $N(0)$ post-softmax weights are equal and have total mass at most 1, thus $\|a(\beta)\|_\infty \leq 1/N(0)$. This gives entropy bound

$$H(\beta) \geq \sum_{j=1}^n a_j(\beta) \log(N(0)) = \log N(0) \geq \log 2,$$

since post-softmax weights sum to 1. The conclusion for the sequence $(s^{(n)})$ follows from the fact that $H(\beta) \geq \log 2$ for infinitely many n . $\square$

Definition 6.2.2: Given a score vector $s \in \mathbb{R}^n$, define the *rank free energy* $\mathcal{F}(\beta) = \log Z(\beta)$ and the *maximum post-softmax gap* $\mathcal{G}(\beta) = \max_{j,k} |a_j(\beta) - a_k(\beta)|$.

Proposition 6.2.3 (Rank-Collapse Criterion): Let $(s^{(n)})$ be a sequence of score vectors such that $s^{(n)} \in \mathbb{R}^n$. For any positive sequence (β_n) , the following are equivalent:

1. $\mathcal{F}(\beta_n) \rightarrow \infty$,
2. $\mathcal{G}(\beta_n) \rightarrow 0$,
3. $\max_j |a_j^{(n)}(\beta_n) - 1/n| \rightarrow 0$.

Proof. For simplicity, we omit the argument β from the post-softmax weights $a_j^{(n)}(\beta)$. The maximum value of s has $\Delta_j = 0$, so

$$\max_j a_j^{(n)} = \frac{1}{Z(\beta)} = e^{-\mathcal{F}(\beta)}.$$

Since $\mathcal{G}(\beta) \leq \max_j a_j^{(n)}$, $\mathcal{F}(\beta_n) \rightarrow \infty$ implies that $\mathcal{G}(\beta_n) \rightarrow 0$.

Conversely, by a basic property of the arithmetic mean, we have the bound $\min_j a_j^{(n)}(\beta) \leq \frac{1}{n}$, which gives

$$\max_j a_j^{(n)} = \max_j a_j^{(n)} - \min_k a_k^{(n)} + \min_k a_k^{(n)} \leq \mathcal{G}(\beta) + \frac{1}{n}.$$

Thus, $\mathcal{G}(\beta_n) \rightarrow 0$ implies that $\max_j a_j^{(n)} \rightarrow 0$, which is equivalent to $\mathcal{F}(\beta) \rightarrow \infty$.

Finally, we observe that $\max_j a_j^{(n)} \rightarrow 0$ is equivalent to $\max_j |a_j^{(n)}(\beta_n) - 1/n| \rightarrow 0$, since $0 \leq a_k^{(n)} \leq \max_j a_j^{(n)}$ and $1/n \rightarrow 0$. $\square$

6.3 Resolved Accumulation

Definition 6.3.1 (Rank Boundary): For $0 \leq r \leq \log n$, define

$$B(r) = \sup\{\beta \geq 0 : \mathcal{F}(\beta) \geq r\} = \sup\{\beta \geq 0 : Z(\beta) \geq e^r\}.$$

Definition 6.3.2 (Resolved Accumulation Scale): For $r \geq 0$ and a fixed $s \in \mathbb{R}^n$, define

$$\Lambda^{(r)} = \sup_{t > 0 : \log N(t) > r} \frac{\log N(t) - r}{t},$$

with convention $\sup \emptyset = 0$.

Proposition 6.3.3: Let (r_n) be a positive sequence such that $r_n \rightarrow \infty$. If $\beta_n < \Lambda_n^{(r_n)}$ eventually, then $\mathcal{G}(\beta_n) \rightarrow 0$.

Proof. By assumption, for large enough n we have $\beta_n < \Lambda_n^{(r_n)}$. Hence, there is some t_n such that

$$\beta_n < \frac{\log N(t_n) - r_n}{t_n},$$

since $\Lambda_n^{(r_n)}$ is defined in terms of a supremum. Rearranging,

$$N(t_n)e^{-\beta_n t_n} > e^{r_n}.$$

Since $N(t_n)$ competitors each contribute at least $e^{-\beta_n t_n}$ to $Z(\beta_n)$, we have that

$$Z(\beta_n) > e^{r_n},$$

which gives that $\mathcal{F}(\beta_n) \rightarrow \infty$, and thus $\mathcal{G}(\beta_n) \rightarrow 0$ by Proposition 6.2.3. $\square$

Proposition 6.3.4: Define the Laplace envelope

$$S(\beta) = \sup_{t > 0} (\log N(t) - \beta t).$$

For every $\beta > 0$,

$$S(\beta) \leq \mathcal{F}(\beta) \leq S(\beta) + \log(1 + \log n).$$

Proof. Fix $t > 0$. The $N(t)$ tokens with $\Delta_j \leq t$ each contribute at least $e^{-\beta t}$, thus

$$N(t)e^{-\beta t} \leq Z(\beta),$$

and hence $\log N(t) - \beta t \leq \mathcal{F}(\beta)$. After taking supremums, this gives $S(\beta) \leq \mathcal{F}(\beta)$.

Consider the ordered gaps $0 = \Delta_1 \leq \Delta_2 \leq \dots \leq \Delta_n$. Since $j \leq N(\Delta_j)$,

$$\log j - \beta \Delta_j \leq \log N(\Delta_j) - \beta \Delta_j \leq S(\beta),$$

by definition of S . Thus,

$$e^{-\beta \Delta_j} \leq \frac{S(\beta)}{j},$$

which when summed over j gives

$$Z(\beta) \leq e^{S(\beta)} \leq S(\beta) \sum_{j=1}^n \frac{1}{j} \leq e^{S(\beta)} (1 + \log n),$$

by a classic bound on the harmonic numbers. Taking logarithms on both sides of the inequality yields the desired upper bound for $\mathcal{F}$. $\square$

Corollary 6.3.1: If $\beta_n/\Lambda_n \rightarrow 0$ and the *contact-count entropy* $C_n = \log N(\Delta_\Lambda^{(n)}) = \Lambda_n \Delta_\Lambda^{(n)}$ satisfies $C_n \rightarrow \infty$, then $\mathcal{G}(\beta_n) \rightarrow 0$.

Proof. Set $r_n = \beta_n/\Lambda_n$. By the previous proposition,

$$Z(\beta_n) \geq N(\Delta_n) e^{-\beta_n \Delta_\Lambda^{(n)}} = e^{C_n - r_n C_n} = e^{(1-r_n)C_n},$$

and thus $\mathcal{F}(\beta_n) \geq (1 - r_n)C_n$. Since $r_n \rightarrow 0$ and $C_n \rightarrow \infty$, we have that $\mathcal{F} \rightarrow \infty$, which is equivalent to $\mathcal{G}(\beta_n) \rightarrow 0$ by Proposition 6.2.3. $\square$

References.

- [Vas+17] A. Vaswani *et al.*, “Attention Is All You Need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. [Online]. Available: <https://arxiv.org/abs/1706.03762>
- [Nak25] K. M. Nakanishi, “Scalable-Softmax is Superior for Attention.” [Online]. Available: <https://arxiv.org/abs/2501.19399>
- [Vel+24] P. Veličković, C. Perivolaropoulos, F. Barbero, and R. Pascanu, “Softmax is not enough (for sharp size generalisation),” *arXiv preprint arXiv:2410.01104*, 2024, [Online]. Available: <https://arxiv.org/abs/2410.01104>
- [Che+25] S. Chen, Z. Lin, Y. Polyanskiy, and P. Rigollet, “Critical attention scaling in long-context transformers,” *arXiv preprint arXiv:2510.05554*, 2025, [Online]. Available: <https://arxiv.org/abs/2510.05554>
- [LLR83] M. R. Leadbetter, G. Lindgren, and H. Rootzén, *Extremes and Related Properties of Random Sequences and Processes*. in Springer Series in Statistics. New York: Springer-Verlag, 1983.
- [GG25] A. Giorlandino and S. Goldt, “Two failure modes of deep transformers and how to avoid them: a unified theory of signal propagation at initialisation.” [Online]. Available: <https://arxiv.org/abs/2505.24333>
- [Ver26] R. Vershynin, *High-Dimensional Probability: An Introduction with Applications in Data Science*, 2nd ed. in Cambridge Series in Statistical and Probabilistic Mathematics. Cambridge: Cambridge University Press, 2026.
- [HK26] T. Hayase and R. Karakida, “A Unified Framework for Critical Scaling of Inverse Temperature in Self-Attention,” *arXiv preprint arXiv:2605.12697*, 2026, [Online]. Available: <https://arxiv.org/abs/2605.12697>